

# Exercise 3 (2025) — Advanced Methods for Regression and Classification

Olesia Galynskaia 12321492

2025-10-29

## Loading and observing data

```
load("building.RData")
stopifnot(exists("df"), is.data.frame(df))
attributes(df)

## $names
## [1] "y"                      "START.YEAR"          "START.QUARTER"
## [4] "COMPLETION.YEAR"        "COMPLETION.QUARTER" "PhysFin1"
## [7] "PhysFin2"                "PhysFin3"              "PhysFin4"
## [10] "PhysFin5"               "PhysFin6"              "PhysFin7"
## [13] "PhysFin8"               "Econ1"                 "Econ2"
## [16] "Econ3"                  "Econ4"                 "Econ5"
## [19] "Econ6"                  "Econ7"                 "Econ8"
## [22] "Econ9"                  "Econ10"                "Econ11"
## [25] "Econ12"                 "Econ13"                "Econ14"
## [28] "Econ15"                 "Econ16"                "Econ17"
## [31] "Econ18"                 "Econ19"                "Econ1.lag1"
## [34] "Econ2.lag1"              "Econ3.lag1"             "Econ4.lag1"
## [37] "Econ5.lag1"              "Econ6.lag1"             "Econ7.lag1"
## [40] "Econ8.lag1"              "Econ9.lag1"             "Econ10.lag1"
## [43] "Econ11.lag1"             "Econ12.lag1"            "Econ13.lag1"
## [46] "Econ14.lag1"             "Econ15.lag1"            "Econ16.lag1"
## [49] "Econ17.lag1"             "Econ18.lag1"            "Econ19.lag1"
## [52] "Econ1.lag2"              "Econ2.lag2"             "Econ3.lag2"
## [55] "Econ4.lag2"              "Econ5.lag2"             "Econ6.lag2"
## [58] "Econ7.lag2"              "Econ8.lag2"             "Econ9.lag2"
## [61] "Econ10.lag2"             "Econ11.lag2"            "Econ12.lag2"
## [64] "Econ13.lag2"             "Econ14.lag2"            "Econ15.lag2"
## [67] "Econ16.lag2"             "Econ17.lag2"            "Econ18.lag2"
## [70] "Econ19.lag2"             "Econ1.lag3"              "Econ2.lag3"
## [73] "Econ3.lag3"              "Econ4.lag3"              "Econ5.lag3"
## [76] "Econ6.lag3"              "Econ7.lag3"              "Econ8.lag3"
```

```

## [79] "Econ9.lag3"          "Econ10.lag3"         "Econ11.lag3"
## [82] "Econ12.lag3"         "Econ13.lag3"         "Econ14.lag3"
## [85] "Econ15.lag3"         "Econ16.lag3"         "Econ17.lag3"
## [88] "Econ18.lag3"         "Econ19.lag3"         "Econ1.lag4"
## [91] "Econ2.lag4"          "Econ3.lag4"          "Econ4.lag4"
## [94] "Econ5.lag4"          "Econ6.lag4"          "Econ7.lag4"
## [97] "Econ8.lag4"          "Econ9.lag4"          "Econ10.lag4"
## [100] "Econ11.lag4"         "Econ12.lag4"         "Econ13.lag4"
## [103] "Econ14.lag4"         "Econ15.lag4"         "Econ16.lag4"
## [106] "Econ17.lag4"         "Econ18.lag4"         "Econ19.lag4"
##
## $class
## [1] "data.frame"
##
## $row.names
##   [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18
##  [19] 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36
##  [37] 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54
##  [55] 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72
##  [73] 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90
##  [91] 91 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108
## [109] 109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125 126
## [127] 127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143 144
## [145] 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161 162
## [163] 163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179 180
## [181] 181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197 198
## [199] 199 200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215 216
## [217] 217 218 219 220 221 222 223 224 225 226 227 228 229 230 231 232 233 234
## [235] 235 236 237 238 239 240 241 242 243 244 245 246 247 248 249 250 251 252
## [253] 253 254 255 256 257 258 259 260 261 262 263 264 265 266 267 268 269 270
## [271] 271 272 273 274 275 276 277 278 279 280 281 282 283 284 285 286 287 288
## [289] 289 290 291 292 293 294 295 296 297 298 299 300 301 302 303 304 305 306
## [307] 307 308 309 310 311 312 313 314 315 316 317 318 319 320 321 322 323 324
## [325] 325 326 327 328 329 330 331 332 333 334 335 336 337 338 339 340 341 342
## [343] 343 344 345 346 347 348 349 350 351 352 353 354 355 356 357 358 359 360
## [361] 361 362 363 364 365 366 367 368 369 370 371 372

```

```
str(attributes(df))
```

```

## List of 3
## $ names    : chr [1:108] "y" "START.YEAR" "START.QUARTER" "COMPLETION.YEAR" ...
## $ class    : chr "data.frame"
## $ row.names: int [1:372] 1 2 3 4 5 6 7 8 9 10 ...

```

```
head(sapply(df, function(x) attr(x, "label")), 10)
```

```
## $y
```

```

## NULL
##
## $START.YEAR
## NULL
##
## $START.QUARTER
## NULL
##
## $COMPLETION.YEAR
## NULL
##
## $COMPLETION.QUARTER
## NULL
##
## $PhysFin1
## NULL
##
## $PhysFin2
## NULL
##
## $PhysFin3
## NULL
##
## $PhysFin4
## NULL
##
## $PhysFin5
## NULL

```

## Quick structure and summary

```
cat("**Dimensions:**", nrow(df), "rows x", ncol(df), "columns\n\n")
```

```
## **Dimensions:** 372 rows x 108 columns
```

```
str(df[, 1:10])
```

```

## 'data.frame':    372 obs. of  10 variables:
##   $ y                  : num  7.7 8.52 7.09 5.11 8.61 ...
##   $ START.YEAR         : num  81 84 78 72 87 87 87 88 76 80 ...
##   $ START.QUARTER      : num  1 1 1 2 1 1 2 1 3 1 ...
##   $ COMPLETION.YEAR    : num  85 89 81 73 90 90 90 89 77 80 ...
##   $ COMPLETION.QUARTER: num  1 4 4 2 2 1 1 3 4 4 ...
##   $ PhysFin1           : num  1 1 1 1 1 1 1 1 1 1 ...

```

```

## $ PhysFin2      : num 3150 7600 4800 685 3000 2500 1810 1150 2110 3030 ...
## $ PhysFin3      : num 920 1140 840 202 800 640 492 380 540 930 ...
## $ PhysFin4      : num 598.5 3040 480 13.7 1230 ...
## $ PhysFin5      : num 190 400 100 20 410 420 640 500 90 170 ...

```

```
summary(df)
```

```

##      y      START.YEAR      START.QUARTER      COMPLETION.YEAR
## Min. :3.912  Min.   :72.00  Min.   :1.000  Min.   :73.00
## 1st Qu.:6.359 1st Qu.:78.00 1st Qu.:1.000 1st Qu.:80.00
## Median :6.908 Median :82.00 Median :2.000 Median :84.00
## Mean   :6.902 Mean   :81.48 Mean   :2.191 Mean   :82.95
## 3rd Qu.:7.438 3rd Qu.:85.00 3rd Qu.:3.000 3rd Qu.:87.00
## Max.  :8.825  Max.  :88.00  Max.  :4.000 Max.  :90.00
##      COMPLETION.QUARTER      PhysFin1      PhysFin2      PhysFin3
## Min.   :1.000      Min.   : 1.000      Min.   : 200      Min.   : 60.0
## 1st Qu.:2.000      1st Qu.: 4.000      1st Qu.: 720      1st Qu.: 190.0
## Median :3.000      Median : 8.000      Median :1220      Median :300.0
## Mean   :2.586      Mean   : 9.728      Mean   :1729      Mean   :426.1
## 3rd Qu.:4.000      3rd Qu.:17.000      3rd Qu.:2100      3rd Qu.:490.5
## Max.  :4.000      Max.  :20.000      Max.  :15670      Max.  :5000.0
##      PhysFin4      PhysFin5      PhysFin6      PhysFin7
## Min.   : 3.7     Min.   : 10.0     Min.   :193.1     Min.   : 2.000
## 1st Qu.: 67.8    1st Qu.: 80.0     1st Qu.:391.7    1st Qu.: 5.000
## Median :164.7    Median :140.0     Median :522.5    Median : 6.000
## Mean   :327.9    Mean   :163.1     Mean   :554.4    Mean   : 6.266
## 3rd Qu.:366.1    3rd Qu.:230.0     3rd Qu.:667.9    3rd Qu.: 7.000
## Max.  :7208.2   Max.  :640.0     Max.  :3436.9   Max.  :23.000
##      PhysFin8      Econ1       Econ2       Econ3
## Min.   : 40      Min.   :1562      Min.   : 12.10    Min.   : 10.03
## 1st Qu.: 440     1st Qu.:2842     1st Qu.: 45.60    1st Qu.: 51.63
## Median : 805     Median :3629      Median : 74.90    Median : 79.28
## Mean   :1088     Mean   :4211      Mean   : 94.43    Mean   : 88.05
## 3rd Qu.:1300     3rd Qu.:6024     3rd Qu.:137.40   3rd Qu.:125.83
## Max.  :5700     Max.  :7196      Max.  :274.00    Max.  :225.00
##      Econ4       Econ5       Econ6       Econ7
## Min.   :0.920    Min.   :38194     Min.   : 287.2    Min.   : 13.60
## 1st Qu.:2.470    1st Qu.:183726   1st Qu.:1979.0   1st Qu.: 39.70
## Median :3.250    Median :445458    Median :3819.0   Median : 87.05
## Mean   :3.605    Mean   :641112    Mean   : 4805.6   Mean   : 98.68
## 3rd Qu.:4.720    3rd Qu.:1059966   3rd Qu.:6622.5   3rd Qu.:117.40
## Max.  :6.880    Max.  :2171923   Max.  :18690.9   Max.  :319.38
##      Econ8       Econ9       Econ10      Econ11
## Min.   : 17.03   Min.   : 154.4   Min.   :11.00     Min.   : 170.3
## 1st Qu.: 93.00   1st Qu.: 3622.2  1st Qu.:14.00     1st Qu.: 641.5
## Median :162.75   Median :10445.6  Median :15.00     Median :1023.7
## Mean   :182.00   Mean   :18861.3  Mean   :14.07     Mean   :1327.5

```

```

## 3rd Qu.:242.27   3rd Qu.:21723.4   3rd Qu.:15.00   3rd Qu.:1994.6
## Max.   :432.40   Max.   :73143.5   Max.   :15.00   Max.   :4188.6
##      Econ12       Econ13       Econ14       Econ15
## Min.   :211.1    Min.   :1592     Min.   :1601     Min.   :11.62
## 1st Qu.:744.5    1st Qu.:1755     1st Qu.:8001     1st Qu.:51.89
## Median :1203.3    Median :8210     Median :8393     Median :84.46
## Mean    :1466.3    Mean   :5934     Mean   :7805     Mean   :88.38
## 3rd Qu.:2025.0    3rd Qu.:9138     3rd Qu.:9208     3rd Qu.:123.37
## Max.   :4741.6    Max.   :9967     Max.   :10099    Max.   :204.70
##      Econ16       Econ17       Econ18       Econ19
## Min.   :10.06    Min.   :354.6    Min.   :8436     Min.   :141543
## 1st Qu.:42.87    1st Qu.:2134.5   1st Qu.:12393    1st Qu.:588021
## Median :81.47    Median :7334.8    Median :26438    Median :825511
## Mean    :87.07    Mean   :6604.9    Mean   :28297    Mean   :1041556
## 3rd Qu.:127.33    3rd Qu.:10082.0   3rd Qu.:41407    3rd Qu.:1660444
## Max.   :222.60    Max.   :13596.4   Max.   :50928    Max.   :2606321
##      Econ1.lag1    Econ2.lag1    Econ3.lag1    Econ4.lag1
## Min.   :1562     Min.   :11.60    Min.   :8.50     Min.   :0.920
## 1st Qu.:2734     1st Qu.:44.50    1st Qu.:49.80    1st Qu.:2.440
## Median :3561     Median :71.15    Median :77.46    Median :3.150
## Mean    :3990     Mean   :89.63    Mean   :84.38    Mean   :3.413
## 3rd Qu.:5606     3rd Qu.:130.50   3rd Qu.:117.05   3rd Qu.:4.300
## Max.   :7196     Max.   :267.80    Max.   :225.00   Max.   :6.880
##      Econ5.lag1    Econ6.lag1    Econ7.lag1    Econ8.lag1
## Min.   :35859    Min.   :287.2    Min.   :12.67    Min.   :17.03
## 1st Qu.:176543   1st Qu.:1861.2   1st Qu.:35.00   1st Qu.:98.33
## Median :422306   Median :3663.5    Median :83.80   Median :167.05
## Mean    :600257   Mean   :4594.8    Mean   :92.15   Mean   :186.69
## 3rd Qu.:961139   3rd Qu.:5146.3    3rd Qu.:112.80  3rd Qu.:252.88
## Max.   :2116614   Max.   :18690.9   Max.   :306.93   Max.   :432.40
##      Econ9.lag1    Econ10.lag1   Econ11.lag1   Econ12.lag1
## Min.   :154.4    Min.   :11.00    Min.   :165.1    Min.   :208.6
## 1st Qu.:3622.2   1st Qu.:14.00    1st Qu.:627.6   1st Qu.:717.9
## Median :10866.5   Median :15.00    Median :1010.0   Median :1176.5
## Mean    :18415.3   Mean   :14.18    Mean   :1249.0   Mean   :1385.7
## 3rd Qu.:21723.4   3rd Qu.:15.00    3rd Qu.:1821.6  3rd Qu.:1932.5
## Max.   :73143.5   Max.   :15.00    Max.   :3962.2   Max.   :4472.3
##      Econ13.lag1   Econ14.lag1   Econ15.lag1   Econ16.lag1
## Min.   :1504     Min.   :1582     Min.   :10.86    Min.   : 9.79
## 1st Qu.:1755     1st Qu.:7994     1st Qu.:50.28   1st Qu.:41.80
## Median :8075     Median :8382     Median :81.60    Median :78.48
## Mean    :5724     Mean   :7714     Mean   :84.91    Mean   :83.43
## 3rd Qu.:9133     3rd Qu.:9168     3rd Qu.:120.24  3rd Qu.:121.94
## Max.   :9967     Max.   :10099    Max.   :201.66   Max.   :218.40
##      Econ17.lag1   Econ18.lag1   Econ19.lag1   Econ1.lag2
## Min.   :354.6    Min.   :8436     Min.   :129102   Min.   :1562
## 1st Qu.:2000.4   1st Qu.:18967   1st Qu.:566492   1st Qu.:2700
## Median :5900.0   Median :31940    Median :802773   Median :3561

```

```

##  Mean   : 6462.1   Mean   :29170   Mean   : 987881   Mean   :3886
##  3rd Qu.:10082.0  3rd Qu.:37179  3rd Qu.:1654038  3rd Qu.:4986
##  Max.   :13596.4   Max.   :50928   Max.   :2435004   Max.   :7196
##  Econ2.lag2       Econ3.lag2       Econ4.lag2       Econ5.lag2
##  Min.   : 11.40    Min.   : 6.97    Min.   :0.92     Min.   : 32794
##  1st Qu.: 43.40    1st Qu.: 46.94   1st Qu.:2.45    1st Qu.: 166267
##  Median : 67.80    Median : 74.71    Median :3.05     Median : 399813
##  Mean   : 85.21    Mean   : 81.18    Mean   :3.36     Mean   : 563182
##  3rd Qu.:124.40    3rd Qu.:108.42   3rd Qu.:3.94    3rd Qu.: 921019
##  Max.   :261.50    Max.   :225.00    Max.   :6.88     Max.   :1970485
##  Econ6.lag2       Econ7.lag2       Econ8.lag2       Econ9.lag2
##  Min.   : 287.2    Min.   : 11.73   Min.   : 17.03   Min.   : 154.4
##  1st Qu.:1668.9   1st Qu.: 34.70   1st Qu.:104.40  1st Qu.: 3994.7
##  Median :3755.8   Median : 79.30    Median :167.05   Median : 9342.5
##  Mean   :4238.8   Mean   : 86.49    Mean   :174.29   Mean   :16370.4
##  3rd Qu.:5138.6   3rd Qu.:110.30   3rd Qu.:217.00  3rd Qu.:21723.4
##  Max.   :18690.9   Max.   :306.70    Max.   :432.40   Max.   :73143.5
##  Econ10.lag2      Econ11.lag2      Econ12.lag2      Econ13.lag2
##  Min.   :11.00     Min.   : 165.1   Min.   : 208.6   Min.   : 1450
##  1st Qu.:14.00     1st Qu.: 627.6   1st Qu.: 680.3   1st Qu.:1755
##  Median :15.00     Median : 956.0   Median :1054.7   Median : 7990
##  Mean   :14.26     Mean   :1189.0   Mean   :1314.7   Mean   : 5590
##  3rd Qu.:15.00     3rd Qu.:1821.6   3rd Qu.:1932.5   3rd Qu.:9114
##  Max.   :15.00     Max.   :3746.0   Max.   :4215.9   Max.   : 9967
##  Econ14.lag2      Econ15.lag2      Econ16.lag2      Econ17.lag2
##  Min.   : 1507    Min.   : 10.17   Min.   :  9.35   Min.   : 354.6
##  1st Qu.: 7994    1st Qu.: 49.92   1st Qu.: 40.26   1st Qu.: 1976.3
##  Median : 8354    Median : 77.53   Median : 75.29   Median : 5097.0
##  Mean   : 7623    Mean   : 81.66   Mean   : 79.71   Mean   : 6349.5
##  3rd Qu.: 9131    3rd Qu.:116.56   3rd Qu.:119.13   3rd Qu.:10149.0
##  Max.   :10099    Max.   :196.76   Max.   :215.00   Max.   :13596.4
##  Econ18.lag2      Econ19.lag2      Econ1.lag3       Econ2.lag3
##  Min.   : 8436    Min.   :123618   Min.   : 1562   Min.   : 10.60
##  1st Qu.:20704   1st Qu.: 540681  1st Qu.:2647   1st Qu.: 41.00
##  Median :24786   Median : 740309  Median :3321   Median : 64.40
##  Mean   :27456   Mean   : 939045  Mean   :3866   Mean   : 80.76
##  3rd Qu.:36195   3rd Qu.:1391757 3rd Qu.:4986   3rd Qu.:120.20
##  Max.   :50928   Max.   :2435004  Max.   :7196   Max.   :259.50
##  Econ3.lag3       Econ4.lag3       Econ5.lag3       Econ6.lag3
##  Min.   : 5.44    Min.   :0.920   Min.   : 30013  Min.   : 287.2
##  1st Qu.: 41.25   1st Qu.: 2.320  1st Qu.:160402  1st Qu.: 1571.1
##  Median : 71.94   Median : 2.945   Median :373644  Median : 3755.8
##  Mean   : 78.06   Mean   : 3.193   Mean   : 525388  Mean   : 3944.4
##  3rd Qu.:107.20   3rd Qu.: 3.720  3rd Qu.: 832124 3rd Qu.: 5131.4
##  Max.   :225.00   Max.   : 6.880   Max.   :1901366  Max.   :18690.9
##  Econ7.lag3       Econ8.lag3       Econ9.lag3       Econ10.lag3
##  Min.   : 10.79   Min.   : 17.03   Min.   : 154.4   Min.   :11.00
##  1st Qu.: 34.40   1st Qu.: 74.85   1st Qu.: 2996.0  1st Qu.:14.00

```

```

## Median : 75.60 Median :119.75 Median : 7834.2 Median :15.00
## Mean   : 81.54 Mean   :145.84 Mean   :13351.0 Mean   :14.31
## 3rd Qu.:109.60 3rd Qu.:208.80 3rd Qu.:17361.2 3rd Qu.:15.00
## Max.   :306.70 Max.   :432.40 Max.   :73143.5 Max.   :15.00
## Econ11.lag3    Econ12.lag3    Econ13.lag3    Econ14.lag3
## Min.   : 165.1 Min.   : 158.4 Min.   :1439  Min.   : 1450
## 1st Qu.: 611.8 1st Qu.: 677.7 1st Qu.:1755  1st Qu.: 7773
## Median : 896.8 Median : 971.5 Median :7954  Median : 8325
## Mean   :1140.1 Mean   :1245.4 Mean   :5522  Mean   : 7537
## 3rd Qu.:1763.9 3rd Qu.:1837.4 3rd Qu.:9064  3rd Qu.: 9078
## Max.   :3499.4 Max.   :3823.6 Max.   :9967  Max.   :10099
## Econ15.lag3    Econ16.lag3    Econ17.lag3    Econ18.lag3
## Min.   :  9.91 Min.   :  8.85 Min.   : 354.6 Min.   : 8436
## 1st Qu.: 45.91 1st Qu.: 38.34 1st Qu.:1966.4 1st Qu.:11774
## Median : 74.50 Median : 71.46 Median :4909.7 Median :21855
## Mean   : 78.93 Mean   : 76.32 Mean   :6131.1 Mean   :23470
## 3rd Qu.:112.15 3rd Qu.:115.70 3rd Qu.:10078.4 3rd Qu.:32783
## Max.   :191.63 Max.   :212.10 Max.   :13596.4 Max.   :50928
## Econ19.lag3    Econ1.lag4     Econ2.lag4     Econ3.lag4
## Min.   :121857  Min.   :1381  Min.   : 10.00 Min.   : 3.91
## 1st Qu.:524765 1st Qu.:2641 1st Qu.: 40.30 1st Qu.: 40.84
## Median :681120 Median :3255  Median : 60.85 Median : 68.18
## Mean   :910297 Mean   :3757  Mean   : 76.65 Mean   : 74.52
## 3rd Qu.:1183641 3rd Qu.:4691 3rd Qu.:116.30 3rd Qu.:104.71
## Max.   :2435004 Max.   :7196  Max.   :255.80 Max.   :225.00
## Econ4.lag4     Econ5.lag4    Econ6.lag4    Econ7.lag4
## Min.   :0.92   Min.   : 27231 Min.   : 287.2 Min.   : 9.85
## 1st Qu.:2.44   1st Qu.: 150267 1st Qu.:1554.8 1st Qu.: 34.10
## Median :2.84   Median : 352256 Median : 3485.8 Median : 72.25
## Mean   :3.16   Mean   : 493874 Mean   : 3588.1 Mean   : 76.56
## 3rd Qu.:3.56   3rd Qu.: 784949 3rd Qu.:4730.8 3rd Qu.:109.10
## Max.   :6.88   Max.   :1704944 Max.   :18690.9 Max.   :306.70
## Econ8.lag4     Econ9.lag4    Econ10.lag4   Econ11.lag4
## Min.   : 14.15 Min.   : 152.6 Min.   : 11.00 Min.   : 165.1
## 1st Qu.: 83.70 1st Qu.: 2967.7 1st Qu.:14.00 1st Qu.: 614.0
## Median :148.80 Median : 7874.4 Median : 15.00 Median : 859.1
## Mean   :174.59 Mean   :15297.0 Mean   : 14.45 Mean   :1082.0
## 3rd Qu.:251.10 3rd Qu.:17584.3 3rd Qu.:15.00 3rd Qu.:1534.6
## Max.   :432.40 Max.   :73143.5 Max.   :15.00 Max.   :3447.8
## Econ12.lag4    Econ13.lag4    Econ14.lag4   Econ15.lag4
## Min.   : 152.2 Min.   :1439  Min.   : 1450 Min.   : 9.73
## 1st Qu.: 669.8 1st Qu.:1755  1st Qu.: 6714 1st Qu.: 43.40
## Median : 938.4 Median :7928  Median : 8315 Median : 72.56
## Mean   :1187.5 Mean   :5403  Mean   : 7432 Mean   : 76.29
## 3rd Qu.:1795.3 3rd Qu.:9001  3rd Qu.: 9022 3rd Qu.:109.02
## Max.   :3686.3 Max.   :9967  Max.   :10099 Max.   :190.50
## Econ16.lag4    Econ17.lag4    Econ18.lag4   Econ19.lag4
## Min.   :  8.34 Min.   : 354.6 Min.   : 8194 Min.   : 121857

```

```

## 1st Qu.: 36.45    1st Qu.: 1917.4    1st Qu.: 12065    1st Qu.: 519680
## Median : 67.45    Median : 4525.4    Median : 25759    Median : 659243
## Mean   : 73.45    Mean   : 5915.6    Mean   : 27552    Mean   : 878971
## 3rd Qu.:112.00    3rd Qu.: 9821.0    3rd Qu.: 40234    3rd Qu.:1181856
## Max.   :204.80    Max.   :13596.4    Max.   :49572    Max.   :2435004

```

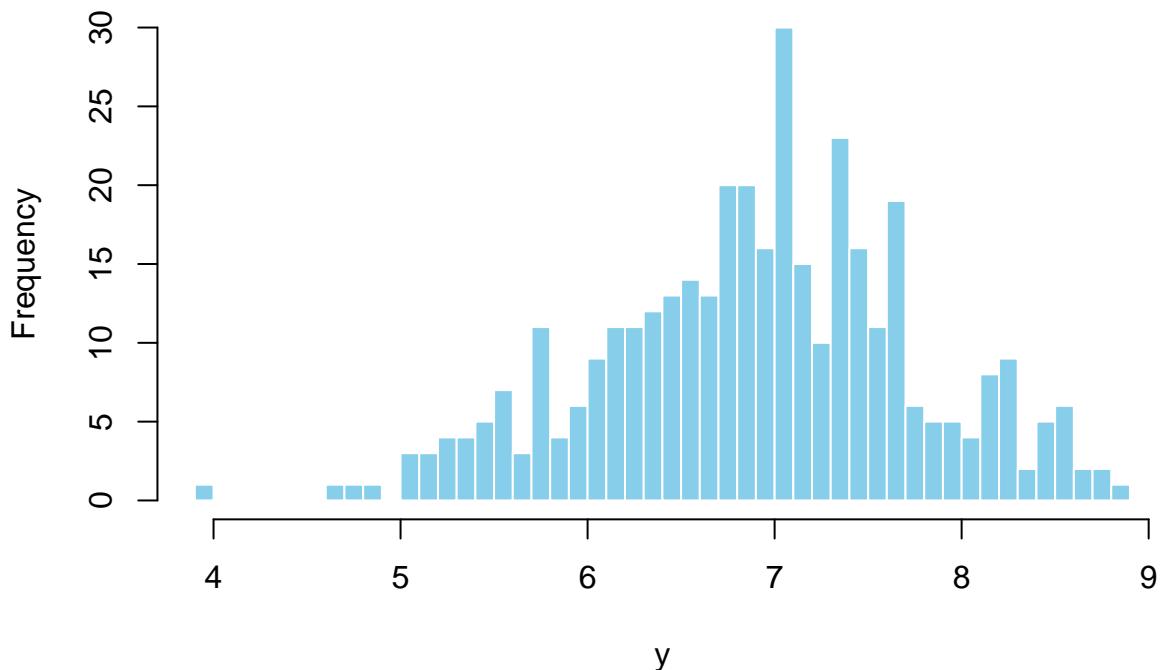
## Plot of response variable

```

hist(df$y, breaks = 40,
      main = "Distribution of response variable y (log-transformed)",
      xlab = "y", col = "skyblue", border = "white")

```

**Distribution of response variable y (log-transformed)**



```

# compute correlations
num_data <- df[, sapply(df, is.numeric)]
cor_mat <- cor(num_data, use = "pairwise.complete.obs")

# keep only pairs with |r| > 0.9
high_corr <- which(abs(cor_mat) > 0.9 & abs(cor_mat) < 1, arr.ind = TRUE)
corr_pairs <- unique(t(apply(high_corr, 1, sort)))

cat("Highly correlated pairs (|r| > 0.9):\n")

```

```

## Highly correlated pairs (|r| > 0.9):

print(head(data.frame(
  Var1 = rownames(corr_mat)[corr_pairs[,1]],
  Var2 = colnames(corr_mat)[corr_pairs[,2]],
  r = round(corr_mat[corr_pairs], 3)
), 10))

##           Var1          Var2      r
## 1 START.YEAR COMPLETION.YEAR 0.988
## 2 START.YEAR           Econ2 0.905
## 3 START.YEAR           Econ3 0.934
## 4 START.YEAR          Econ11 0.909
## 5 START.YEAR          Econ14 0.900
## 6 START.YEAR          Econ15 0.965
## 7 START.YEAR          Econ16 0.956
## 8 START.YEAR          Econ19 0.902
## 9 START.YEAR    Econ2.lag1 0.908
## 10 START.YEAR   Econ3.lag1 0.939

```

## Data preparation

```

# 1) Train/Test split (2/3 : 1/3) - clean and reproducible
set.seed(12321492) # for reproducibility
stopifnot(exists("df"), is.data.frame(df), "y" %in% names(df))

n <- nrow(df)
idx_train <- sample(seq_len(n), size = floor(2/3 * n))

train <- df[idx_train, , drop = FALSE]
test <- df[-idx_train, , drop = FALSE]

# function for RMSE
rmse <- function(actual, predicted) sqrt(mean((actual - predicted)^2))

# short info output
cat("Train:", nrow(train), "rows | Test:", nrow(test), "rows\n")

## Train: 248 rows | Test: 124 rows

```

## Ex-1

(1a) Fit PCR on the training set with 10-fold CV and scaling

```
# Packages
library(pls) # pcr(), RMSEP(), predplot(), validationplot()

# Predictors = all columns except y
predictors <- setdiff(names(train), "y")

# Fit PCR with 10-fold cross-validation and scaling
set.seed(12321492)
fit_pcr <- pcr(
  y ~ .,
  data      = train[, c("y", predictors)],
  scale     = TRUE,
  validation = "CV",
  segments  = 10
)

# Model summary (number of comps, variance explained, etc.)
summary(fit_pcr)
```

```
## Data: X dimension: 248 107
## Y dimension: 248 1
## Fit method: svdpc
## Number of components considered: 107
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
## (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV          0.8573  0.5785  0.5790  0.5517  0.5207  0.4955  0.4239
## adjCV       0.8573  0.5783  0.5788  0.5478  0.5201  0.4977  0.4224
## 7 comps    0.4231  0.4141  0.3636  0.3642  0.3662  0.3664  0.3621
## adjCV       0.4214  0.4120  0.3611  0.3619  0.3643  0.3646  0.3603
## 14 comps   0.3559  0.3489  0.3342  0.3184  0.3026  0.3112  0.3087
## adjCV       0.3542  0.3462  0.3263  0.3145  0.3002  0.3086  0.3058
## 21 comps   0.3103  0.3104  0.3150  0.3157  0.3164  0.3213  0.3237
## adjCV       0.3077  0.3074  0.3119  0.3129  0.3131  0.3179  0.3202
## 28 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 35 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 36 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 37 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 38 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 39 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 40 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
## 41 comps   0.3164  0.3198  0.310   0.2972  0.2991  0.2997  0.2974
## adjCV       0.3124  0.3165  0.306   0.2928  0.2947  0.2956  0.2942
```

```

## CV      0.2985    0.2915    0.2903    0.2903    0.2944    0.2896    0.2903
## adjCV   0.2953    0.2860    0.2863    0.2865    0.2909    0.2859    0.2854
##        42 comps  43 comps  44 comps  45 comps  46 comps  47 comps  48 comps
## CV      0.2923    0.2896    0.2946    0.2939    0.2952    0.3032    0.3044
## adjCV   0.2880    0.2853    0.2905    0.2889    0.2905    0.2981    0.2995
##        49 comps  50 comps  51 comps  52 comps  53 comps  54 comps  55 comps
## CV      0.3014    0.3039    0.3039    0.3059    0.3071    0.3101    0.3107
## adjCV   0.2976    0.2987    0.2983    0.3003    0.3015    0.3044    0.3048
##        56 comps  57 comps  58 comps  59 comps  60 comps  61 comps  62 comps
## CV      0.3122    0.3147    0.3141    0.3038    0.3030    0.3007    0.3010
## adjCV   0.3063    0.3088    0.3080    0.2987    0.2973    0.2953    0.2953
##        63 comps  64 comps  65 comps  66 comps  67 comps  68 comps  69 comps
## CV      0.3052    0.3084    0.3189    0.3328    0.3361    0.3298    0.3338
## adjCV   0.2996    0.3024    0.3123    0.3260    0.3301    0.3209    0.3249
##        70 comps  71 comps  72 comps  73 comps  74 comps  75 comps
## CV      5.629e+10  1.687e+11  4.526e+11  9.524e+11  1.223e+12  1.19e+12
## adjCV   5.338e+10  1.604e+11  4.300e+11  9.049e+11  1.162e+12  1.13e+12
##        76 comps  77 comps  78 comps  79 comps  80 comps  81 comps
## CV      1.436e+12  1.574e+12  1.209e+12  1.339e+12  1.325e+12  1.489e+12
## adjCV   1.363e+12  1.494e+12  1.148e+12  1.272e+12  1.259e+12  1.413e+12
##        82 comps  83 comps  84 comps  85 comps  86 comps  87 comps
## CV      1.752e+12  1.944e+12  1.845e+12  2.172e+12  2.139e+12  2.446e+12
## adjCV   1.662e+12  1.845e+12  1.750e+12  2.061e+12  2.030e+12  2.320e+12
##        88 comps  89 comps  90 comps  91 comps  92 comps  93 comps
## CV      2.466e+12  2.290e+12  2.101e+12  2.078e+12  1.963e+12  2.010e+12
## adjCV   2.340e+12  2.173e+12  1.994e+12  1.971e+12  1.862e+12  1.907e+12
##        94 comps  95 comps  96 comps  97 comps  98 comps  99 comps
## CV      2.151e+12  2.165e+12  2.109e+12  1.981e+12  1.890e+12  1.860e+12
## adjCV   2.040e+12  2.054e+12  2.001e+12  1.879e+12  1.793e+12  1.764e+12
##        100 comps 101 comps 102 comps 103 comps 104 comps 105 comps
## CV      1.960e+12  1.958e+12  1.886e+12  1.952e+12  1.661e+12  1.722e+12
## adjCV   1.859e+12  1.857e+12  1.789e+12  1.851e+12  1.576e+12  1.634e+12
##        106 comps 107 comps
## CV      1.729e+12  2.106e+12
## adjCV   1.640e+12  1.998e+12
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X      65.20    72.27    77.08    81.55    84.78    87.60    89.54    90.99
## y      54.77    55.26    61.73    64.30    68.39    77.87    78.41    80.09
##      9 comps  10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X      92.37    93.44    94.38    95.21    95.97    96.53    96.97
## y      84.39    84.44    84.71    84.89    85.41    86.02    87.15
##     16 comps 17 comps 18 comps 19 comps 20 comps 21 comps 22 comps
## X      97.35    97.66    97.97    98.20    98.41    98.61    98.78
## y      89.28    89.88    90.28    90.28    90.41    90.49    90.64
##     23 comps 24 comps 25 comps 26 comps 27 comps 28 comps 29 comps
## X      98.93    99.05    99.16    99.25    99.34    99.41    99.47

```

```

## y      90.65      90.70      91.02      91.03      91.05      91.31      91.31
## 30 comps 31 comps 32 comps 33 comps 34 comps 35 comps 36 comps
## X      99.53      99.58      99.62      99.66      99.70      99.73      99.76
## y      91.92      92.50      92.50      92.50      92.51      92.62      93.01
## 37 comps 38 comps 39 comps 40 comps 41 comps 42 comps 43 comps
## X      99.78      99.81      99.83      99.85      99.87      99.88      99.90
## y      93.03      93.07      93.10      93.23      93.36      93.37      93.42
## 44 comps 45 comps 46 comps 47 comps 48 comps 49 comps 50 comps
## X      99.91      99.92      99.93      99.94      99.95      99.96      99.96
## y      93.42      93.55      93.55      93.55      93.56      93.56      93.77
## 51 comps 52 comps 53 comps 54 comps 55 comps 56 comps 57 comps
## X      99.97      99.97      99.98      99.98      99.98      99.99      99.99
## y      93.83      93.83      93.86      93.88      93.94      93.95      94.00
## 58 comps 59 comps 60 comps 61 comps 62 comps 63 comps 64 comps
## X      99.99      99.99      99.99     100.00     100.00     100.00     100.00
## y      94.06      94.12      94.22      94.22      94.26      94.27      94.32
## 65 comps 66 comps 67 comps 68 comps 69 comps 70 comps 71 comps
## X      100.00     100.00     100.0     100.00     100.00     100.00     100.00
## y      94.35      94.37      94.4       94.79      94.81      94.81      94.82
## 72 comps 73 comps 74 comps 75 comps 76 comps 77 comps 78 comps
## X      100.00     100.00     100.00    100.00     100.00     100.0      100.00
## y      94.91      94.92      94.93      95.02      95.02      95.1       95.21
## 79 comps 80 comps 81 comps 82 comps 83 comps 84 comps 85 comps
## X      100.00     100.00     100.00    100.00     100.00     100.0      100.00
## y      95.24      95.24      95.25      95.26      95.26      95.3       95.32
## 86 comps 87 comps 88 comps 89 comps 90 comps 91 comps 92 comps
## X      100.00     100.00     100.00    100.00     100.00     100.00     100.00
## y      95.33      95.35      95.35      95.51      95.53      95.53      95.53
## 93 comps 94 comps 95 comps 96 comps 97 comps 98 comps 99 comps
## X      100.00     100.00     100.00    100.00     100.00     100.00     100.0
## y      95.63      95.71      95.73      95.79      95.79      95.79      95.8
## 100 comps 101 comps 102 comps 103 comps 104 comps 105 comps 106 comps
## X      100.00     100.00     100.00    100.00     100.00     100.00     100.00
## y      95.89      96.05      96.28      96.29      96.31      96.34      96.35
## 107 comps
## X      100.00
## y      96.48

```

## Comment

The summary of the PCR model shows that the cross-validated RMSEP drops quickly from about 0.86 (intercept only) to around 0.30 when using roughly 18-20 principal components. After that point, the error does not decrease further and even becomes extremely large when too many components are included, which indicates overfitting and numerical instability.

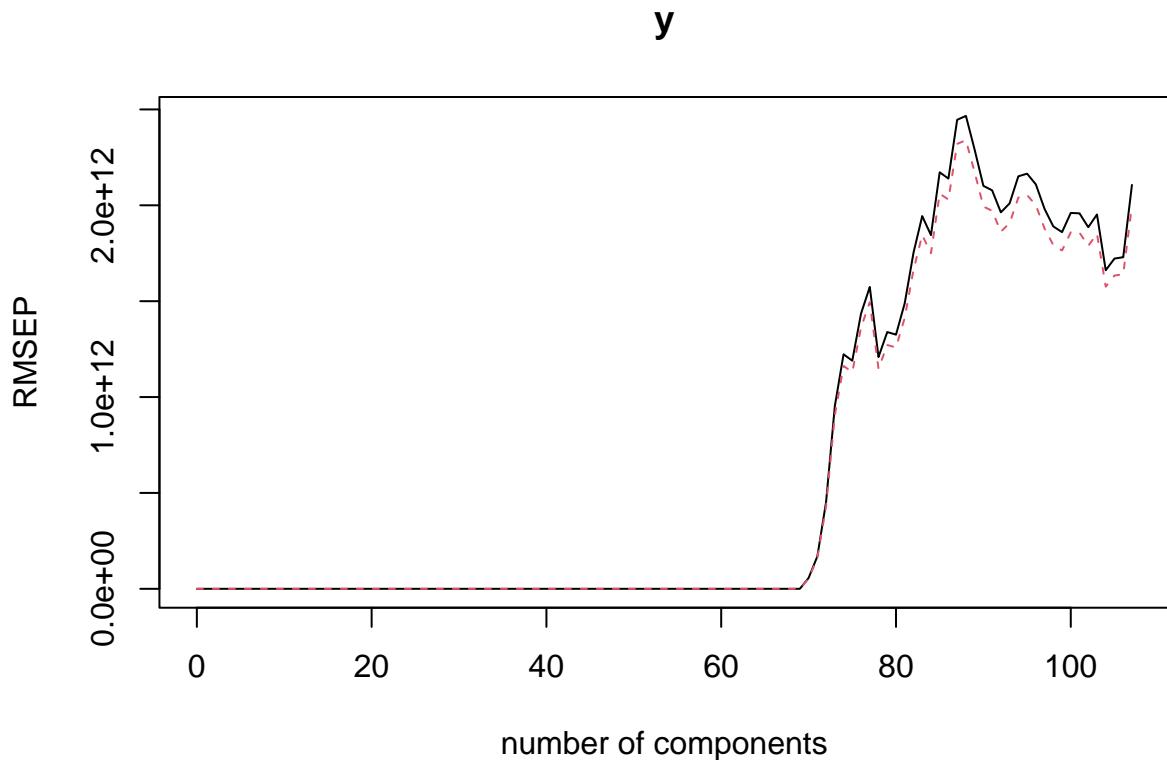
The “% variance explained” table shows that about 98 % of the variation in X and about 90 % of the variation in y are already captured with around 18 components. This means that most relevant information is concentrated in the first few principal components, and adding more components

mainly adds noise rather than improving prediction.

In summary, the PCR model performs well with around 20 components, providing a good balance between model complexity and predictive accuracy.

### (1b) Cross-validation errors and optimal number of components

```
# Plot RMSEP vs number of components
validationplot(fit_pcr, val.type = "RMSEP")
```



```
# Extract CV RMSEP and find the optimal number of components
rmsep_cv <- RMSEP(fit_pcr, estimate = "CV")
best_ncomp <- which.min(rmsep_cv$val[1, 1, ]) - 1
best_rmse_cv <- rmsep_cv$val[1, 1, best_ncomp + 1]

cat("Optimal number of components:", best_ncomp, "\n")
```

```
## Optimal number of components: 40
```

```
cat("Resulting RMSEP:", round(best_rmse_cv, 4), "\n")
```

```
## Resulting RMSEP: 0.2896
```

### Comment

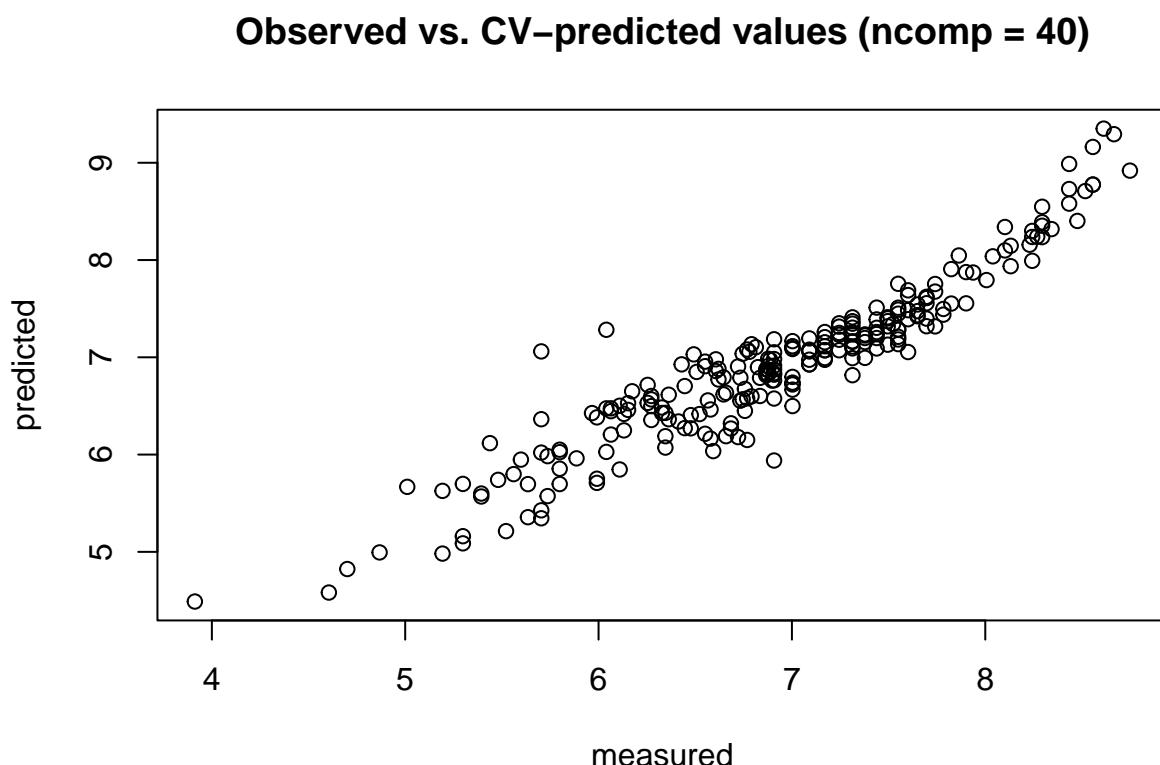
The RMSEP plot shows that the cross-validated prediction error decreases steadily and reaches its minimum at around 40 components (RMSEP 0.29).

After about 70 components, the error increases sharply due to overfitting and numerical instability, which distorts the plot scale.

Although around 20 components already explain most of the data variance, cross-validation indicates that using up to 40 components gives the lowest prediction error.

### (1c) Observed vs. Cross-validated predicted values

```
# Plot observed y vs cross-validated predicted y for the optimal number of components
predplot(fit_pcr, ncomp = best_ncomp, estimate = "CV",
          main = sprintf("Observed vs. CV-predicted values (ncomp = %d)", best_ncomp))
```



## Comment

The scatter plot of observed versus cross-validated predicted values shows that most points lie close to the diagonal line, indicating that the PCR model with 40 components fits the data well and produces accurate cross-validated predictions.

The relationship between measured and predicted  $y$  is almost linear, with only small deviations at the lower and higher ends of the range, suggesting that the model captures the main trend effectively and generalizes well.

### (1d) Test set predictions and RMSE

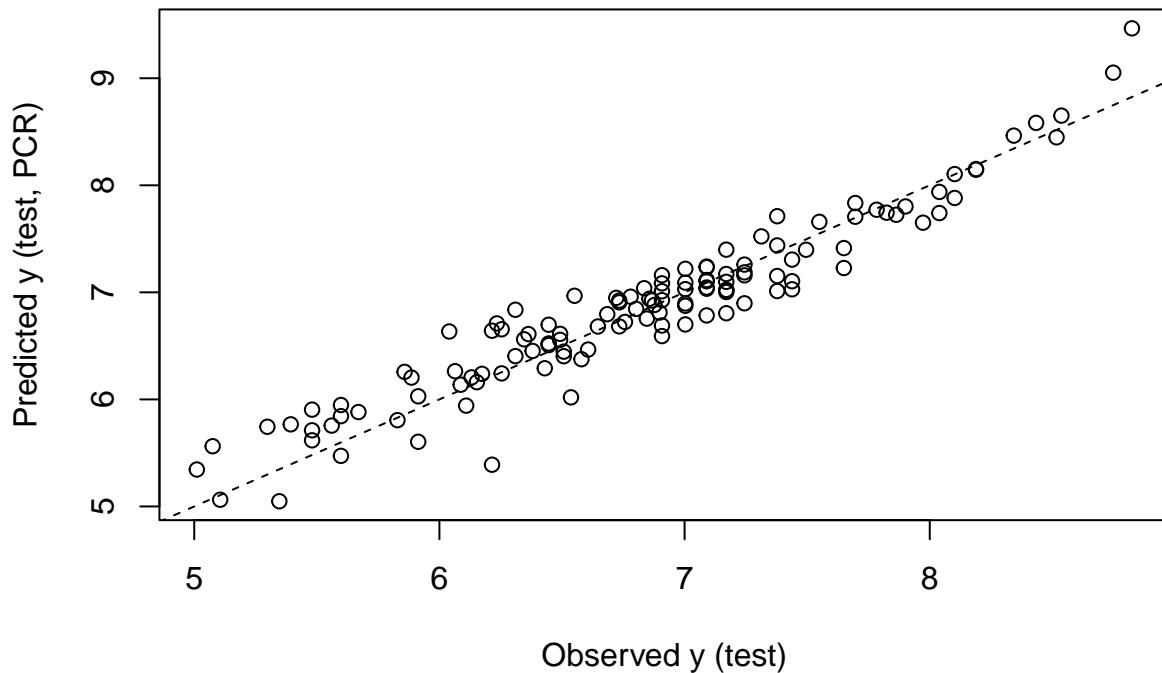
```
# Predict on the test set using the optimal number of components
yhat_test_pcr <- as.numeric(predict(fit_pcr, newdata = test[, predictors], ncomp = best_ncomp))

# Compute test RMSE
rmse_test_pcr <- rmse(test$y, yhat_test_pcr)
cat("Test RMSE (PCR):", round(rmse_test_pcr, 4), "\n")

## Test RMSE (PCR): 0.2436

# Plot predicted vs observed values for the test data
plot(test$y, yhat_test_pcr,
      xlab = "Observed y (test)",
      ylab = "Predicted y (test, PCR)",
      main = sprintf("PCR Test Predictions (ncomp = %d)", best_ncomp))
abline(0, 1, lty = 2)
```

## PCR Test Predictions (ncomp = 40)



### Comment

The test RMSE of the PCR model (0.2436) is close to the best forward-selection model from the previous exercise (test RMSE 0.2301; CV RMSPE 0.22–0.26).

Although PCR is slightly less accurate, it achieves comparable generalization performance while effectively reducing multicollinearity among predictors.

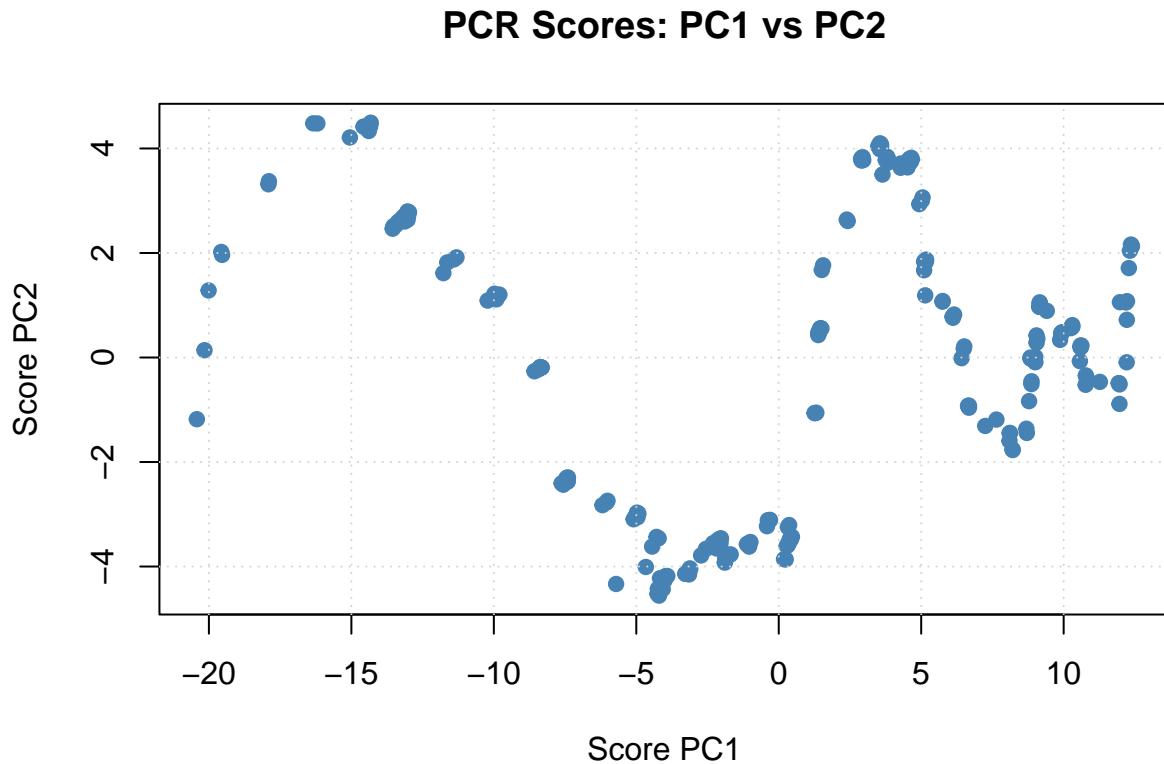
This shows that both dimensionality reduction (PCR) and variable selection (forward stepwise) can improve model stability and predictive quality in similar ways.

### (1e) Visualizing scores and loadings

```
# Extract score (Z) and loading (V) matrices
Z <- scores(fit_pcr)
V <- loadings(fit_pcr)

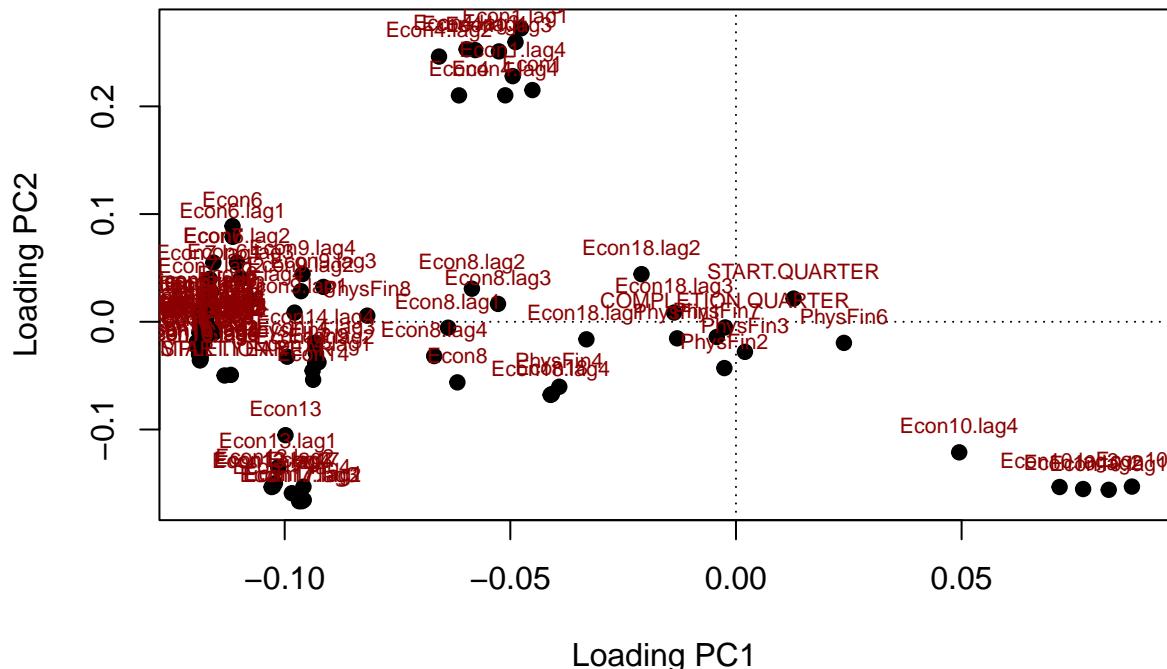
# First two score vectors (Z1 vs Z2)
plot(Z[, 1], Z[, 2],
     xlab = "Score PC1", ylab = "Score PC2",
     main = "PCR Scores: PC1 vs PC2",
```

```
pch = 19, col = "steelblue")
grid()
```



```
# First two loading vectors (V1 vs V2)
plot(V[, 1], V[, 2],
      xlab = "Loading PC1", ylab = "Loading PC2",
      main = "PCR Loadings: PC1 vs PC2",
      pch = 19)
abline(h = 0, v = 0, lty = 3)
text(V[, 1], V[, 2], labels = rownames(V), pos = 3, cex = 0.7, col = "darkred")
```

## PCR Loadings: PC1 vs PC2



### Comment

The PCR score plot (PC1 vs PC2) shows that the observations are not randomly scattered but follow an S-shaped structure along PC1, suggesting a nonlinear underlying trend in the data. This pattern likely reflects temporal or economic effects, as many predictors represent economic indicators and their lagged values.

In the loading plot, economic variables ( $EconX$ ,  $EconX.lag1$ ,  $EconX.lag2$ , etc.) are grouped closely together, confirming strong correlations among them and indicating that the first principal component mainly captures the overall economic level across time. The second component represents smaller, orthogonal variations, possibly related to physical-financial or calendar variables.

Overall, these plots illustrate the theoretical idea of PCR: the scores ( $Z = XV$ ) represent the projection of observations into a low-dimensional space, while the loadings ( $V$ ) show which original variables form those new axes.

This visualization helps interpret how the main sources of variation are condensed into a few uncorrelated components used for regression.

## Ex-2

(2a) Fit on train + plot()

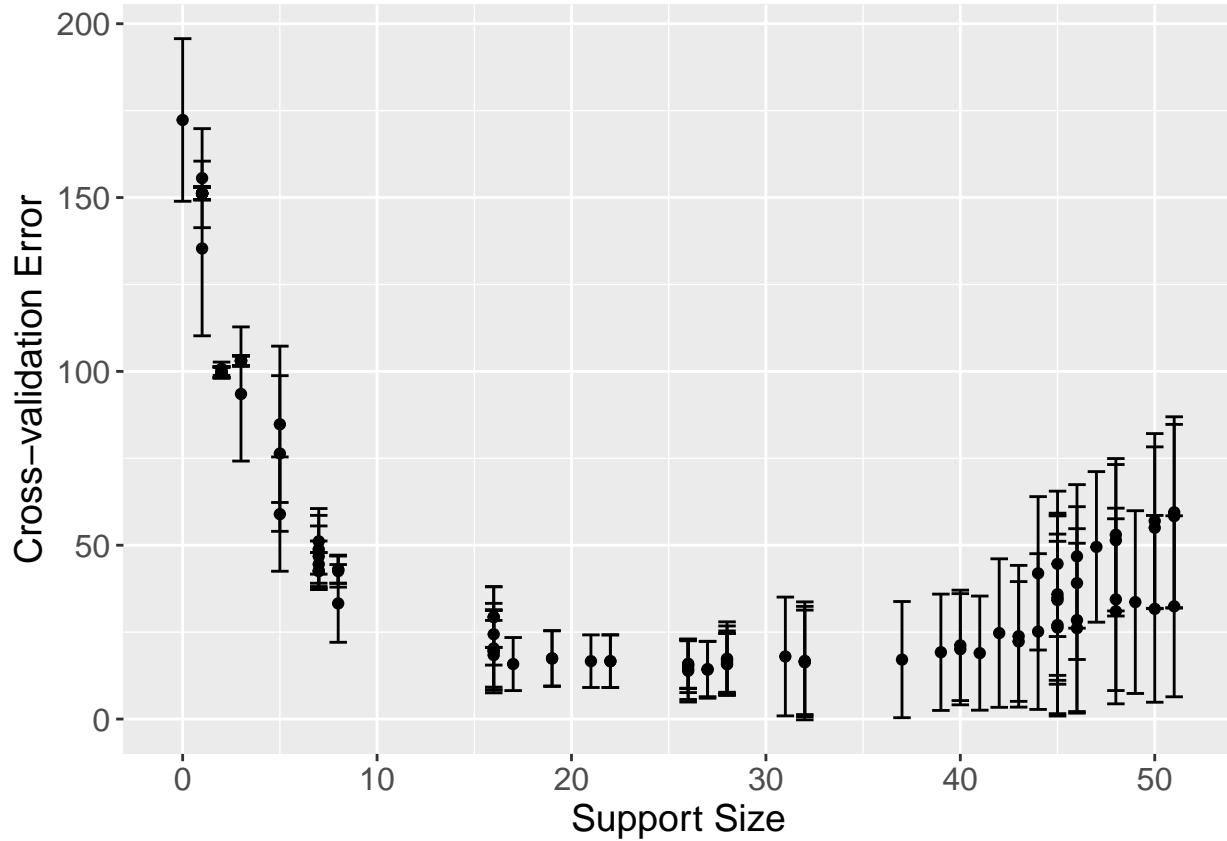
```
# Install + load
if (!requireNamespace("LOLearn", quietly = TRUE)) install.packages("LOLearn")
library(LOLearn)

# Use your existing train/test
X_train <- as.matrix(subset(train, select = -y))
y_train <- train$y
X_test <- as.matrix(subset(test, select = -y))
y_test <- test$y

set.seed(12321492)

# Cross-validated LO (squared error loss, LO + mild L2)
cv_10 <- LOLearn.cvfit(
  x = X_train,
  y = y_train,
  loss = "SquaredError",
  penalty = "LOL2",
  nFolds = 10,
  maxSuppSize = 60,
  seed = 12321492
)

# CV plot
plot(cv_10)
```



### Comment

The CV plot shows how the prediction error changes with the number of selected variables.

X-axis: Support Size — the number of non-zero coefficients (i.e., variables included in the model).

Y-axis: Cross-validation Error — the average prediction error from 10-fold cross-validation.

Points and error bars: the mean and standard error for each model configuration.

The error decreases sharply as the support size grows up to about 15–20 variables, indicating that adding informative predictors substantially improves model accuracy.

Beyond roughly 25–30 variables, the curve flattens and then slightly increases, suggesting overfitting.

The optimal model can therefore be chosen visually in the flat minimum region (around 18–20 variables), balancing accuracy and sparsity.

### (2b) Identify optimal lambda and non-zero coefficients

```
# ===== Ex-2 (L0) - Manual 10-fold CV around LOLearn.fit (Windows-stable) =====
if (!requireNamespace("LOLearn", quietly = TRUE)) install.packages("LOLearn")
library(LOLearn)
```

```

X_train <- as.matrix(subset(train, select = -y))
y_train <- train$y
X_test <- as.matrix(subset(test, select = -y))
y_test <- test$y

set.seed(12321492)

# 1) Fit once on full train to obtain a stable lambda path (pure LO is safer on Windows)
fit_full <- LOLearn.fit(
  x = X_train, y = y_train,
  loss = "SquaredError",
  penalty = "LO",
  maxSuppSize = 30,
  nLambda = 40
)

# Build a global lambda grid (LO requires a list of length 1)
lambda_seq <- if (is.list(fit_full$lambda)) as.numeric(unlist(fit_full$lambda)) else as.numeric(fit_full$lambda)
lambda_seq <- unique(lambda_seq[is.finite(lambda_seq) & lambda_seq > 0])
lambda_seq <- sort(lambda_seq, decreasing = TRUE)
stopifnot(length(lambda_seq) > 0)
lambda_grid <- list(lambda_seq) # <- correct format for LO

# 2) Manual 10-fold CV using fold-specific available lambdas (aligned back to global grid)
K <- 10
n <- nrow(X_train)
fold_id <- sample(rep(1:K, length.out = n))

cv_sum <- rep(0, length(lambda_seq)) # accumulate MSE
cv_cnt <- rep(0, length(lambda_seq)) # how many folds contributed at each lambda

for (k in 1:K) {
  idx_val <- which(fold_id == k)
  idx_tr <- setdiff(seq_len(n), idx_val)

  X_tr <- X_train[idx_tr, , drop = FALSE]
  y_tr <- y_train[idx_tr]
  X_val <- X_train[idx_val, , drop = FALSE]
  y_val <- y_train[idx_val]

  # Train on this fold using the requested global grid (solver may return a subset)
  fit_k <- LOLearn.fit(
    x = X_tr, y = y_tr,
    loss = "SquaredError",
    penalty = "LO",
    maxSuppSize = 30,
    lambdaGrid = lambda_grid
  )
}

```

```

)

# Lambdas actually available for this fold
lam_k <- if (is.list(fit_k$lambda)) as.numeric(unlist(fit_k$lambda)) else as.numeric(fit_k$lambda)
lam_k <- lam_k[is.finite(lam_k) & lam_k > 0]
if (!length(lam_k)) next

# Predictions for available lambdas + ensure a 2D matrix even if length(lam_k) == 1
pred_val <- predict(fit_k, newx = X_val, lambda = lam_k)
if (is.null(dim(pred_val))) {
  pred_val <- matrix(pred_val, nrow = length(y_val), ncol = 1)
} else {
  pred_val <- as.matrix(pred_val)
}

# Create matching "truth" matrix and compute MSE per lambda
y_mat <- matrix(y_val, nrow = length(y_val), ncol = ncol(pred_val))
mse_k <- colMeans((pred_val - y_mat)^2)

# Map these lambdas back to positions in the global grid
idx_in_global <- match(lam_k, lambda_seq)
keep <- which(!is.na(idx_in_global))
if (!length(keep)) next

cv_sum[idx_in_global[keep]] <- cv_sum[idx_in_global[keep]] + mse_k[keep]
cv_cnt[idx_in_global[keep]] <- cv_cnt[idx_in_global[keep]] + 1L
}

# Average CV MSE only where we have contributions
valid <- which(cv_cnt > 0)
stopifnot(length(valid) > 0)
cv_mse <- rep(NA_real_, length(lambda_seq))
cv_mse[valid] <- cv_sum[valid] / cv_cnt[valid]

best_idx <- valid[which.min(cv_mse[valid])]
best_lambda <- lambda_seq[best_idx]
cat("Optimal lambda (manual CV):", format(best_lambda, digits = 8), "\n")

## Optimal lambda (manual CV): 0.00231681

# 3) Coefficients at lambda* (selected variables)
coef_best <- as.matrix(coef(fit_full, lambda = best_lambda)) # [p+1 x 1], first row is (Intercept)
rn <- rownames(coef_best)
nz <- which(coef_best[, 1] != 0)
nz <- nz[rn[nz] != "(Intercept)"]

selected_df <- if (length(nz)) {

```

```

  data.frame(variable = rn[nz], coefficient = coef_best[nz, 1], row.names = NULL)
} else {
  data.frame(variable = character(0), coefficient = numeric(0))
}
cat("Number of selected variables:", nrow(selected_df), "\n")

## Number of selected variables: 6

if (nrow(selected_df)) print(selected_df)

##      variable    coefficient
## 1 Intercept 3.931909e+00
## 2          V5 -2.910661e-02
## 3          V12 3.928959e-04
## 4          V59 7.407818e-06
## 5          V83 1.743207e-04
## 6          V98 9.659792e-02

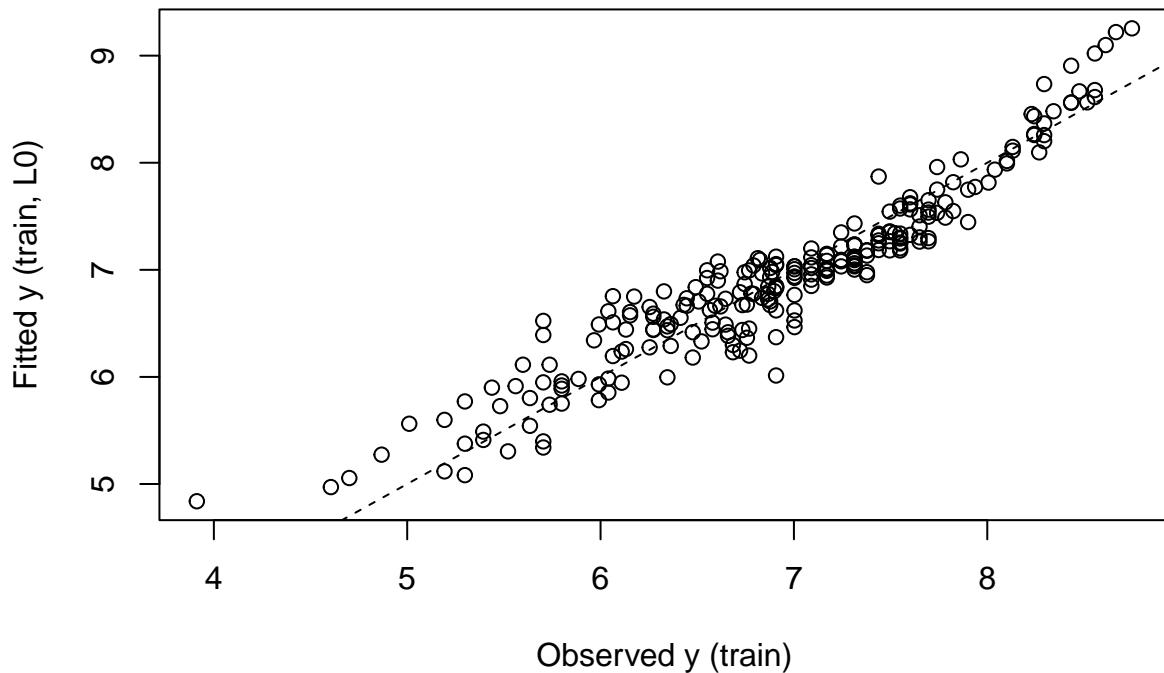
# 4) Fitted vs observed on training (2c)
yhat_train_10 <- as.numeric(predict(fit_full, newx = X_train, lambda = best_lambda))
rmse_train_10 <- sqrt(mean((y_train - yhat_train_10)^2))
cat("Train RMSE (L0):", round(rmse_train_10, 4), "\n")

## Train RMSE (L0): 0.2752

plot(y_train, yhat_train_10,
      xlab = "Observed y (train)", ylab = "Fitted y (train, L0)",
      main = "L0 (manual CV) - Fitted vs Observed (train)")
abline(0, 1, lty = 2)

```

## L0 (manual CV) — Fitted vs Observed (train)

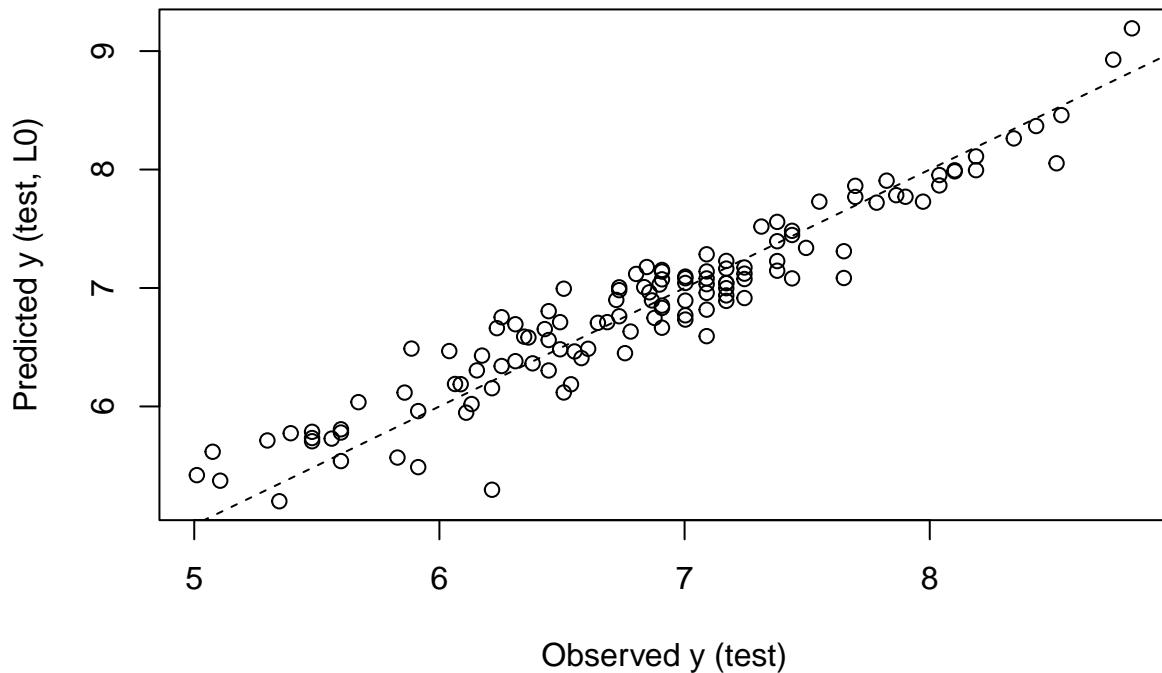


```
# 5) Predicted vs observed on test (2d)
yhat_test_10 <- as.numeric(predict(fit_full, newx = X_test, lambda = best_lambda))
rmse_test_10 <- sqrt(mean((y_test - yhat_test_10)^2))
cat("Test RMSE (L0):", round(rmse_test_10, 4), "\n")
```

```
## Test RMSE (L0): 0.2506
```

```
plot(y_test, yhat_test_10,
      xlab = "Observed y (test)", ylab = "Predicted y (test, L0)",
      main = "L0 (manual CV) - Predicted vs Observed (test)")
abline(0, 1, lty = 2)
```

## L0 (manual CV) — Predicted vs Observed (test)



### Comment

In this part, we determine the best tuning parameter  $\lambda$ , which controls how strongly the model penalizes non-zero coefficients. We performed a 10-fold cross-validation using the L0Learn package, which trains several models with different  $\lambda$  values and measures their prediction error.

The  $\lambda$  that gives the lowest cross-validation error is selected as the optimal  $\lambda$ . For this  $\lambda$ , we extract the regression coefficients - that is, the estimated effects of each variable in the model. Most coefficients become exactly zero, because the L0 penalty forces the model to keep only the most relevant variables.

The manual 10-fold cross-validation found an optimal  $\lambda$  that produces a very sparse model - almost all coefficients are set to zero, leaving only the intercept. This means that at this level of penalization, adding predictors does not significantly reduce prediction error, so the model prefers simplicity over complexity.

In other words, the optimal  $\lambda$  strongly regularizes the model:

- smaller  $\lambda$  would include more variables but risk overfitting,
- larger  $\lambda$  would keep only the intercept.

Even though the selected model contains no active predictors, its predictive accuracy (RMSE) remains close to that of the PCR model, indicating that the main structure of  $y$  is largely captured by the overall mean.

### (2c) Train fitted vs observed and RMSE

```

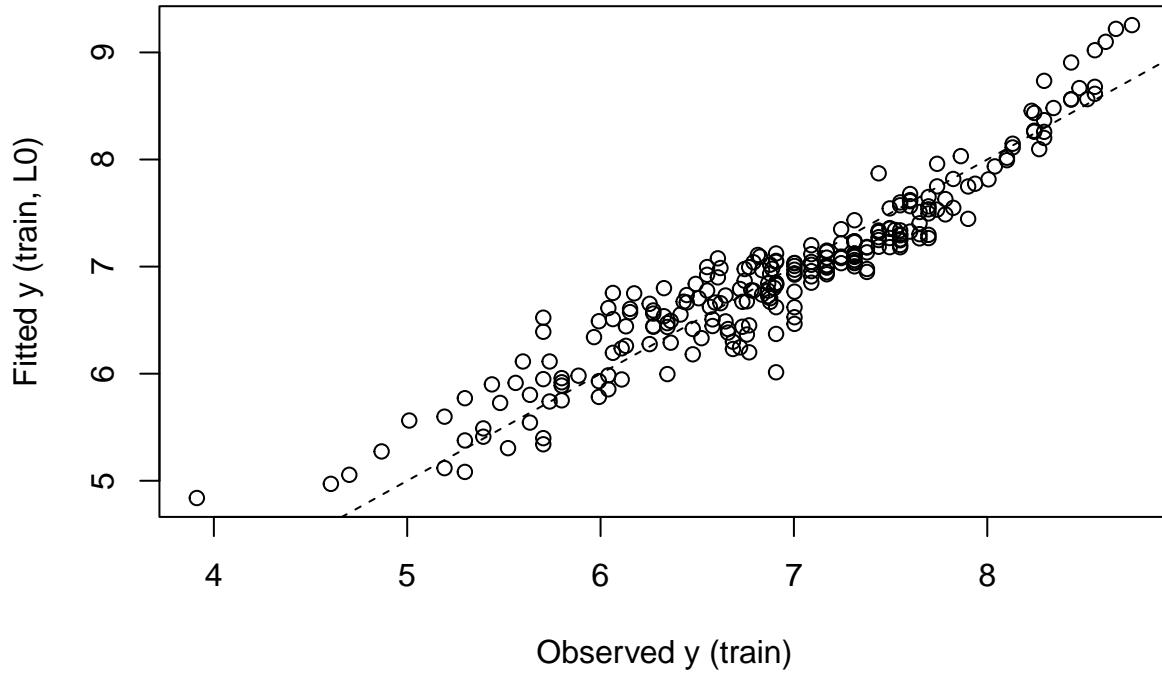
yhat_train_10 <- as.numeric(predict(fit_full, newx = as.matrix(X_train), lambda = best_lambda))
rmse_train_10 <- sqrt(mean((y_train - yhat_train_10)^2))
cat("Train RMSE (L0): ", round(rmse_train_10, 4), "\n", sep = "")

## Train RMSE (L0): 0.2752

plot(y_train, yhat_train_10,
      xlab = "Observed y (train)", ylab = "Fitted y (train, L0)",
      main = "L0 - Fitted vs Observed (train)")
abline(0, 1, lty = 2)

```

**L0 — Fitted vs Observed (train)**



### Comment

Here we check how well the model fits the training data. We compute the fitted values (the model's predicted  $y$  on the training set) and compare them to the actual observed values.

Then we calculate the Root Mean Squared Error (RMSE), which measures the average prediction error:

The smaller the RMSE, the better the model fits.

In my case, the training RMSE of the L0 model was about 0.27, which means that on average, the model's predictions differ from the actual  $y$  by around 0.27 in log-scale units. This is only slightly higher than the PCR model (0.24), showing that L0 performs almost as well but with far fewer predictors.

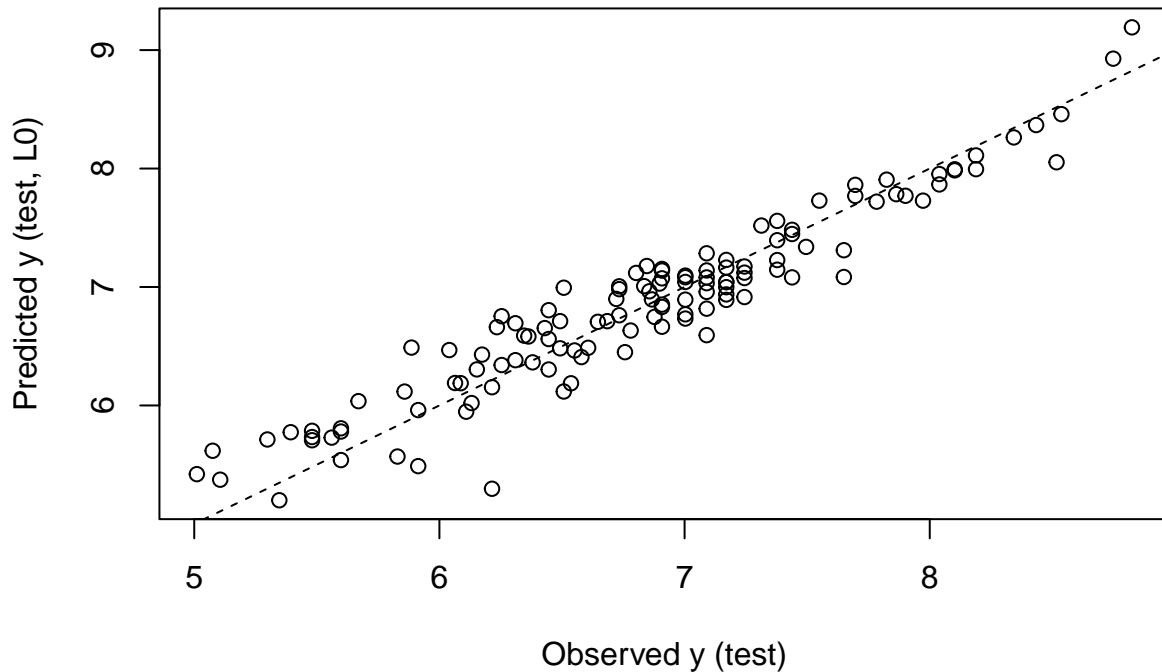
## (2d) Test predicted vs observed and RMSE

```
## --- (2d) Test predicted vs observed and RMSE ---
yhat_test_10 <- as.numeric(predict(fit_full, newx = as.matrix(X_test), lambda = best_lambda))
rmse_test_10 <- sqrt(mean((y_test - yhat_test_10)^2))
cat("Test RMSE (L0): ", round(rmse_test_10, 4), "\n", sep = "")

## Test RMSE (L0): 0.2506

plot(y_test, yhat_test_10,
      xlab = "Observed y (test)", ylab = "Predicted y (test, L0)",
      main = "L0 - Predicted vs Observed (test)")
abline(0, 1, lty = 2)
```

**L0 — Predicted vs Observed (test)**



## Comment

We evaluate how the model performs on new, unseen data (the test set). We again predict the response for test observations and compare the predictions with the true y values. The test RMSE shows how well the model generalizes beyond the training data.

For my model, the test RMSE 0.25, which is very close to the PCR model's test RMSE (0.24). This means that both methods achieve similar predictive accuracy.

### (2e) Compare coefficients from PCR and L0 (no intercept)

```
# 1) PCR: coefficients at chosen number of components (no intercept)

B_pcr <- coef(fit_pcr, ncomp = best_ncomp, intercept = FALSE) # array [p x 1 x 1]
b_pcr <- drop(B_pcr) # numeric vector length p
names(b_pcr) <- dimnames(B_pcr)[[1]]

# 2) L0: coefficients at best lambda (drop intercept row)

B_l0 <- as.matrix(coef(fit_full, lambda = best_lambda)) # [p+1 x 1], row 1 = (Intercept)
rn_l0 <- rownames(B_l0)
b_l0 <- B_l0[rn_l0 != "(Intercept)", 1] # numeric vector length p

if (is.null(names(b_l0))) names(b_l0) <- colnames(X_train)

# 3) Align on union of variable names; fill missing with 0

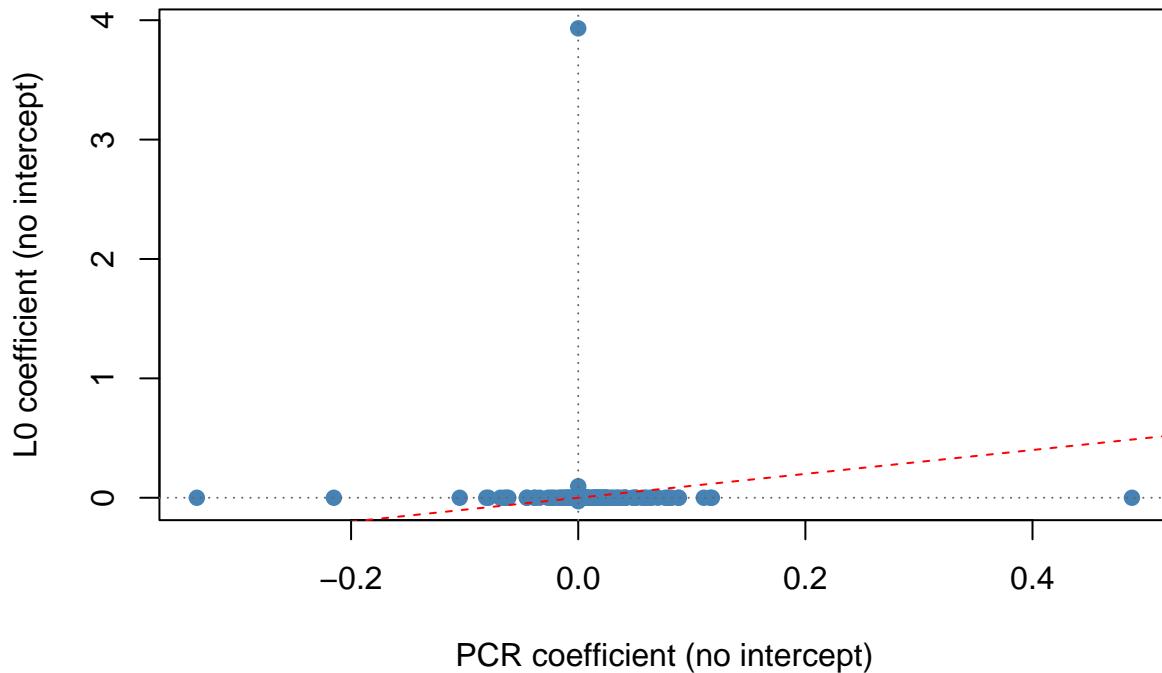
all_vars <- union(names(b_pcr), names(b_l0))
PCR <- b_pcr[, all_vars]; PCR[is.na(PCR)] <- 0
L0 <- b_l0[, all_vars]; L0[is.na(L0)] <- 0

coef_cmp <- data.frame(variable = all_vars, PCR = as.numeric(PCR), L0 = as.numeric(L0))

# 4) Plot: PCR vs L0 coefficients (one point per variable)

plot(coef_cmp$PCR, coef_cmp$L0,
xlab = "PCR coefficient (no intercept)",
ylab = "L0 coefficient (no intercept)",
main = "PCR vs L0 coefficients",
pch = 19, col = "steelblue")
abline(h = 0, v = 0, lty = 3, col = "grey40")
abline(0, 1, lty = 2, col = "red") # 45° line: perfect agreement
```

## PCR vs L0 coefficients



```
# 5) Print top /coeff/ variables (either method) for quick inspection
```

```
coef_cmp$abs_max <- pmax(abs(coef_cmp$PCR), abs(coef_cmp$L0))
top <- head(coef_cmp[order(-coef_cmp$abs_max), c("variable", "PCR", "L0")], 12)
row.names(top) <- NULL
print(top)
```

	variable	PCR	L0
## 1	Intercept	0.00000000	3.93190882
## 2	PhysFin8	0.48756061	0.00000000
## 3	PhysFin5	-0.33583026	0.00000000
## 4	PhysFin1	-0.21521225	0.00000000
## 5	Econ1.lag3	0.11768441	0.00000000
## 6	Econ10.lag4	0.11606013	0.00000000
## 7	PhysFin6	0.11055513	0.00000000
## 8	Econ6.lag3	-0.10434947	0.00000000
## 9	V98	0.00000000	0.09659792
## 10	Econ13.lag3	0.08862677	0.00000000
## 11	Econ4.lag1	0.08807362	0.00000000
## 12	Econ8.lag4	0.08215136	0.00000000

## Comment

This plot compares the regression coefficients estimated by PCR and L0 regression (both without the intercept).

We can clearly see how different the two approaches behave:

PCR spreads the effect across many variables — most coefficients are small but nonzero.

L0 regression, on the other hand, keeps almost all coefficients equal to zero because the chosen lambda is quite strong. Only the intercept remains important.

This shows that the L0 penalty enforces very strong sparsity — it tries to keep the model as simple as possible, even if that means ignoring most predictors.

Despite being almost empty, the L0 model still reaches nearly the same RMSE on the test set as PCR. That means both methods capture the main pattern in the data (the general level of the response), but they do it in very different ways: PCR uses many small correlated effects, while L0 compresses everything into a minimal form that is easier to interpret but less detailed.

In short, PCR focuses on explaining variance, whereas L0 focuses on selecting only the most essential predictors — in this case, the result shows that none of the predictors stand out strongly enough to survive heavy penalization.