

Exercise 4 (2025) — Advanced Methods for Regression and Classification

Olesia Galynskaia 12321492

2025-11-05

Loading and observing data

```
load("building.RData")
stopifnot(exists("df"), is.data.frame(df))
attributes(df)

## $names
## [1] "y"                      "START.YEAR"           "START.QUARTER"
## [4] "COMPLETION.YEAR"        "COMPLETION.QUARTER" "PhysFin1"
## [7] "PhysFin2"                "PhysFin3"              "PhysFin4"
## [10] "PhysFin5"               "PhysFin6"              "PhysFin7"
## [13] "PhysFin8"               "Econ1"                 "Econ2"
## [16] "Econ3"                  "Econ4"                 "Econ5"
## [19] "Econ6"                  "Econ7"                 "Econ8"
## [22] "Econ9"                  "Econ10"                "Econ11"
## [25] "Econ12"                 "Econ13"                "Econ14"
## [28] "Econ15"                 "Econ16"                "Econ17"
## [31] "Econ18"                 "Econ19"                "Econ1.lag1"
## [34] "Econ2.lag1"              "Econ3.lag1"             "Econ4.lag1"
## [37] "Econ5.lag1"              "Econ6.lag1"             "Econ7.lag1"
## [40] "Econ8.lag1"              "Econ9.lag1"             "Econ10.lag1"
## [43] "Econ11.lag1"             "Econ12.lag1"            "Econ13.lag1"
## [46] "Econ14.lag1"             "Econ15.lag1"            "Econ16.lag1"
## [49] "Econ17.lag1"             "Econ18.lag1"            "Econ19.lag1"
## [52] "Econ1.lag2"              "Econ2.lag2"             "Econ3.lag2"
## [55] "Econ4.lag2"              "Econ5.lag2"             "Econ6.lag2"
## [58] "Econ7.lag2"              "Econ8.lag2"             "Econ9.lag2"
## [61] "Econ10.lag2"             "Econ11.lag2"            "Econ12.lag2"
## [64] "Econ13.lag2"             "Econ14.lag2"            "Econ15.lag2"
## [67] "Econ16.lag2"             "Econ17.lag2"            "Econ18.lag2"
## [70] "Econ19.lag2"             "Econ1.lag3"              "Econ2.lag3"
## [73] "Econ3.lag3"              "Econ4.lag3"              "Econ5.lag3"
## [76] "Econ6.lag3"              "Econ7.lag3"              "Econ8.lag3"
```

```

## [79] "Econ9.lag3"          "Econ10.lag3"         "Econ11.lag3"
## [82] "Econ12.lag3"         "Econ13.lag3"         "Econ14.lag3"
## [85] "Econ15.lag3"         "Econ16.lag3"         "Econ17.lag3"
## [88] "Econ18.lag3"         "Econ19.lag3"         "Econ1.lag4"
## [91] "Econ2.lag4"          "Econ3.lag4"          "Econ4.lag4"
## [94] "Econ5.lag4"          "Econ6.lag4"          "Econ7.lag4"
## [97] "Econ8.lag4"          "Econ9.lag4"          "Econ10.lag4"
## [100] "Econ11.lag4"         "Econ12.lag4"         "Econ13.lag4"
## [103] "Econ14.lag4"         "Econ15.lag4"         "Econ16.lag4"
## [106] "Econ17.lag4"         "Econ18.lag4"         "Econ19.lag4"
##
## $class
## [1] "data.frame"
##
## $row.names
##   [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18
##  [19] 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36
##  [37] 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54
##  [55] 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72
##  [73] 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90
##  [91] 91 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108
## [109] 109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125 126
## [127] 127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143 144
## [145] 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161 162
## [163] 163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179 180
## [181] 181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197 198
## [199] 199 200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215 216
## [217] 217 218 219 220 221 222 223 224 225 226 227 228 229 230 231 232 233 234
## [235] 235 236 237 238 239 240 241 242 243 244 245 246 247 248 249 250 251 252
## [253] 253 254 255 256 257 258 259 260 261 262 263 264 265 266 267 268 269 270
## [271] 271 272 273 274 275 276 277 278 279 280 281 282 283 284 285 286 287 288
## [289] 289 290 291 292 293 294 295 296 297 298 299 300 301 302 303 304 305 306
## [307] 307 308 309 310 311 312 313 314 315 316 317 318 319 320 321 322 323 324
## [325] 325 326 327 328 329 330 331 332 333 334 335 336 337 338 339 340 341 342
## [343] 343 344 345 346 347 348 349 350 351 352 353 354 355 356 357 358 359 360
## [361] 361 362 363 364 365 366 367 368 369 370 371 372

```

```
str(attributes(df))
```

```

## List of 3
## $ names    : chr [1:108] "y" "START.YEAR" "START.QUARTER" "COMPLETION.YEAR" ...
## $ class    : chr "data.frame"
## $ row.names: int [1:372] 1 2 3 4 5 6 7 8 9 10 ...

```

```
head(sapply(df, function(x) attr(x, "label")), 100)
```

```
## $y
```

```
## NULL
##
## $START.YEAR
## NULL
##
## $START.QUARTER
## NULL
##
## $COMPLETION.YEAR
## NULL
##
## $COMPLETION.QUARTER
## NULL
##
## $PhysFin1
## NULL
##
## $PhysFin2
## NULL
##
## $PhysFin3
## NULL
##
## $PhysFin4
## NULL
##
## $PhysFin5
## NULL
##
## $PhysFin6
## NULL
##
## $PhysFin7
## NULL
##
## $PhysFin8
## NULL
##
## $Econ1
## NULL
##
## $Econ2
## NULL
##
## $Econ3
## NULL
##
## $Econ4
```

```
## NULL
##
## $Econ5
## NULL
##
## $Econ6
## NULL
##
## $Econ7
## NULL
##
## $Econ8
## NULL
##
## $Econ9
## NULL
##
## $Econ10
## NULL
##
## $Econ11
## NULL
##
## $Econ12
## NULL
##
## $Econ13
## NULL
##
## $Econ14
## NULL
##
## $Econ15
## NULL
##
## $Econ16
## NULL
##
## $Econ17
## NULL
##
## $Econ18
## NULL
##
## $Econ19
## NULL
##
## $Econ1.lag1
```

```
## NULL
##
## $Econ2.lag1
## NULL
##
## $Econ3.lag1
## NULL
##
## $Econ4.lag1
## NULL
##
## $Econ5.lag1
## NULL
##
## $Econ6.lag1
## NULL
##
## $Econ7.lag1
## NULL
##
## $Econ8.lag1
## NULL
##
## $Econ9.lag1
## NULL
##
## $Econ10.lag1
## NULL
##
## $Econ11.lag1
## NULL
##
## $Econ12.lag1
## NULL
##
## $Econ13.lag1
## NULL
##
## $Econ14.lag1
## NULL
##
## $Econ15.lag1
## NULL
##
## $Econ16.lag1
## NULL
##
## $Econ17.lag1
```

```
## NULL
##
## $Econ18.lag1
## NULL
##
## $Econ19.lag1
## NULL
##
## $Econ1.lag2
## NULL
##
## $Econ2.lag2
## NULL
##
## $Econ3.lag2
## NULL
##
## $Econ4.lag2
## NULL
##
## $Econ5.lag2
## NULL
##
## $Econ6.lag2
## NULL
##
## $Econ7.lag2
## NULL
##
## $Econ8.lag2
## NULL
##
## $Econ9.lag2
## NULL
##
## $Econ10.lag2
## NULL
##
## $Econ11.lag2
## NULL
##
## $Econ12.lag2
## NULL
##
## $Econ13.lag2
## NULL
##
## $Econ14.lag2
```

```
## NULL
##
## $Econ15.lag2
## NULL
##
## $Econ16.lag2
## NULL
##
## $Econ17.lag2
## NULL
##
## $Econ18.lag2
## NULL
##
## $Econ19.lag2
## NULL
##
## $Econ1.lag3
## NULL
##
## $Econ2.lag3
## NULL
##
## $Econ3.lag3
## NULL
##
## $Econ4.lag3
## NULL
##
## $Econ5.lag3
## NULL
##
## $Econ6.lag3
## NULL
##
## $Econ7.lag3
## NULL
##
## $Econ8.lag3
## NULL
##
## $Econ9.lag3
## NULL
##
## $Econ10.lag3
## NULL
##
## $Econ11.lag3
```

```
## NULL
##
## $Econ12.lag3
## NULL
##
## $Econ13.lag3
## NULL
##
## $Econ14.lag3
## NULL
##
## $Econ15.lag3
## NULL
##
## $Econ16.lag3
## NULL
##
## $Econ17.lag3
## NULL
##
## $Econ18.lag3
## NULL
##
## $Econ19.lag3
## NULL
##
## $Econ1.lag4
## NULL
##
## $Econ2.lag4
## NULL
##
## $Econ3.lag4
## NULL
##
## $Econ4.lag4
## NULL
##
## $Econ5.lag4
## NULL
##
## $Econ6.lag4
## NULL
##
## $Econ7.lag4
## NULL
##
## $Econ8.lag4
```

```

## NULL
##
## $Econ9.lag4
## NULL
##
## $Econ10.lag4
## NULL
##
## $Econ11.lag4
## NULL

```

Ex-1

Train/test split

```

df <- df[order(1:nrow(df)), ]    # already in time order, this is a no-op

n <- nrow(df)
cut <- floor(0.8 * n)

train <- df[1:cut, ]
test  <- df[(cut+1):n, ]

```

Comment

If we split the data randomly, we would mix future observations into the training set. That would give the model information about the future that it would not have in reality. In other words, the model would “cheat” and we would get over-optimistic results.

To avoid this leakage, we use a time-based split:

the earlier 80% of the observations are used for training

the last 20% are used as the test set

This way, the model only learns from the past and we evaluate it on the future, which matches the real-world forecasting scenario.

Ex-2

(a) PLS with 10-fold CV, time-aware segments

```

library(pls)

set.seed(12321492)  # reproducibility for CV folds

```

```

# 1) Keep only numeric variables; y must be numeric
is_num <- vapply(train, is.numeric, logical(1))
stopifnot("y" %in% names(train), is_num["y"])
train_num <- train[, is_num, drop = FALSE]

# 2) Drop zero-variance columns (defensive)
nzv <- vapply(train_num, function(x) sd(x, na.rm = TRUE) > 0, logical(1))
train_num <- train_num[, nzv, drop = FALSE]

# 3) Fit PLS with consecutive segments to respect time
pls_fit <- plsr(
  y ~ .,
  data      = train_num,
  scale     = TRUE,
  validation = "CV",
  segments   = 10,
  segment.type = "consecutive"
)

# 4) CV diagnostics
summary(pls_fit)

```

```

## Data: X dimension: 297 107
## Y dimension: 297 1
## Fit method: kernelppls
## Number of components considered: 107
##
## VALIDATION: RMSEP
## Cross-validated using 10 consecutive segments.
##          (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV          0.8735  0.5668  0.3908  0.3381  0.3144  0.2845  0.2823
## adjCV       0.8735  0.5653  0.3892  0.3362  0.3115  0.2816  0.2790
##          7 comps 8 comps 9 comps 10 comps 11 comps 12 comps 13 comps
## CV          0.2859  0.3018  0.3042  0.3041  0.3062  0.3068  0.3053
## adjCV       0.2827  0.2975  0.2993  0.2987  0.3007  0.3012  0.2996
##          14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV          0.3076  0.3104  0.3139  0.3157  0.3182  0.3194  0.3214
## adjCV       0.3017  0.3044  0.3075  0.3092  0.3113  0.3125  0.3143
##          21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV          0.3218  0.3217  0.3193  0.3188  0.3164  0.3154  0.3165
## adjCV       0.3146  0.3145  0.3121  0.3116  0.3094  0.3085  0.3095
##          28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV          0.3180  0.3184  0.3171  0.3167  0.3179  0.3174  0.3183
## adjCV       0.3108  0.3112  0.3101  0.3096  0.3107  0.3102  0.3110
##          35 comps 36 comps 37 comps 38 comps 39 comps 40 comps 41 comps
## CV          0.3174  0.3182  0.3199  0.3209  0.3224  0.3219  0.3218
## adjCV       0.3102  0.3109  0.3125  0.3134  0.3147  0.3143  0.3142

```

```

##          42 comps 43 comps 44 comps 45 comps 46 comps 47 comps 48 comps
## CV      0.3214  0.3213  0.3217  0.3215  0.3221  0.3218  0.3213
## adjCV   0.3138  0.3138  0.3141  0.3139  0.3145  0.3143  0.3138
##          49 comps 50 comps 51 comps 52 comps 53 comps 54 comps 55 comps
## CV      0.3214  0.3216  0.3217  0.3217  0.3216  0.3212  0.3216
## adjCV   0.3139  0.3140  0.3141  0.3141  0.3141  0.3136  0.3140
##          56 comps 57 comps 58 comps 59 comps 60 comps 61 comps 62 comps
## CV      0.3218  0.3218  0.3214  0.3213  0.3210  0.3213  0.3211
## adjCV   0.3142  0.3142  0.3139  0.3137  0.3135  0.3137  0.3135
##          63 comps 64 comps 65 comps 66 comps 67 comps 68 comps 69 comps
## CV      0.3208  0.3208  0.3207  0.3207  0.3208  0.3211  0.3211
## adjCV   0.3133  0.3132  0.3132  0.3132  0.3133  0.3135  0.3136
##          70 comps 71 comps 72 comps 73 comps 74 comps 75 comps 76 comps
## CV      0.3211  0.3212  8886609 8895189 12354349 12354386 12354375
## adjCV   0.3135  0.3136  8425845 8433980 11714012 11714047 11714037
##          77 comps 78 comps 79 comps 80 comps 81 comps 82 comps 83 comps
## CV      12354381 12354374 12354381 12354375 12354378 12354376 12354378
## adjCV   11714043 11714036 11714043 11714038 11714040 11714038 11714040
##          84 comps 85 comps 86 comps 87 comps 88 comps 89 comps 90 comps
## CV      12354380 12354382 12354379 12354376 12354381 12354379 12354375
## adjCV   11714042 11714043 11714041 11714038 11714043 11714041 11714037
##          91 comps 92 comps 93 comps 94 comps 95 comps 96 comps 97 comps
## CV      12354373 12354377 12354377 12354377 12354377 12354375 12354387
## adjCV   11714036 11714039 11714039 11714039 11714040 11714037 11714049
##          98 comps 99 comps 100 comps 101 comps 102 comps 103 comps
## CV      12354385 12354380 12354381 12354381 12354387 12354384
## adjCV   11714046 11714042 11714043 11714043 11714049 11714046
##          104 comps 105 comps 106 comps 107 comps
## CV      12354382 12354387 12354374 12354381
## adjCV   11714044 11714049 11714036 11714043
##
## TRAINING: % variance explained
##          1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X       64.92  69.77  73.87  78.03  80.63  81.88  84.96  87.65
## y       62.28  83.21  88.00  90.60  92.43  93.20  93.40  93.56
##          9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X       89.56  91.05  91.98  93.00  93.59  94.17  94.75
## y       93.77  93.94  94.05  94.15  94.31  94.36  94.40
##          16 comps 17 comps 18 comps 19 comps 20 comps 21 comps 22 comps
## X       95.23  95.97  96.33  97.00  97.48  97.80  98.01
## y       94.44  94.48  94.56  94.58  94.61  94.66  94.72
##          23 comps 24 comps 25 comps 26 comps 27 comps 28 comps 29 comps
## X       98.14  98.31  98.47  98.65  98.80  98.92  99.02
## y       94.79  94.83  94.86  94.89  94.93  94.97  95.00
##          30 comps 31 comps 32 comps 33 comps 34 comps 35 comps 36 comps
## X       99.17  99.23  99.29  99.38  99.44  99.48  99.52
## y       95.01  95.04  95.06  95.08  95.10  95.11  95.12
##          37 comps 38 comps 39 comps 40 comps 41 comps 42 comps 43 comps

```

```

## X    99.56    99.63    99.66    99.67    99.70    99.72    99.74
## y    95.13    95.14    95.15    95.16    95.17    95.17    95.18
## 44 comps 45 comps 46 comps 47 comps 48 comps 49 comps 50 comps
## X    99.77    99.79    99.81    99.83    99.85    99.87    99.88
## y    95.18    95.18    95.18    95.19    95.19    95.19    95.19
## 51 comps 52 comps 53 comps 54 comps 55 comps 56 comps 57 comps
## X    99.89    99.91    99.92    99.94    99.95    99.95    99.96
## y    95.19    95.20    95.20    95.20    95.20    95.20    95.20
## 58 comps 59 comps 60 comps 61 comps 62 comps 63 comps 64 comps
## X    99.97    99.97    99.98    99.98    99.98    99.99    99.99
## y    95.20    95.20    95.20    95.20    95.20    95.20    95.20
## 65 comps 66 comps 67 comps 68 comps 69 comps 70 comps 71 comps
## X    99.99    99.99    100.0     100.0     100.0     100.0     100.0
## y    95.20    95.20    95.2     95.2     95.2     95.2     95.2
## 72 comps 73 comps 74 comps 75 comps 76 comps 77 comps 78 comps
## X    100.0     100.0     100.0     100.0     100.0     100.0     100.0
## y    95.2     95.2     95.2     95.2     95.2     95.2     95.2
## 79 comps 80 comps 81 comps 82 comps 83 comps 84 comps 85 comps
## X    100.0     100.0     100.0     100.0     100.0     100.0     100.0
## y    95.2     95.2     95.2     95.2     95.2     95.2     95.2
## 86 comps 87 comps 88 comps 89 comps 90 comps 91 comps 92 comps
## X    100.0     100.0     100.0     100.0     100.0     100.0     100.0
## y    95.2     95.2     95.2     95.2     95.2     95.2     95.2
## 93 comps 94 comps 95 comps 96 comps 97 comps 98 comps 99 comps
## X    100.0     100.0     100.0     100.0     100.0     100.0     100.0
## y    95.2     95.2     95.2     95.2     95.2     95.2     95.2
## 100 comps 101 comps 102 comps 103 comps 104 comps 105 comps 106 comps
## X    100.0     100.0     100.0     100.0     100.0     100.0     100.0
## y    95.2     95.2     95.2     95.2     95.2     95.2     95.2
## 107 comps
## X    100.0
## y    95.2

```

```

cv_res <- RMSEP(pls_fit)      # CV errors for all components (incl. intercept)
cv_mat <- cv_res$val          # 3D array: [metric, response, component]

# Robust extraction: use indices, not names (names may differ or be NULL)
# - 1st dim index 1 = "CV" RMSE
# - 2nd dim index 1 = first/only response
# - 3rd dim = components (slot 1 is intercept-only)
rmse_all <- drop(cv_mat[1, 1, ])      # includes intercept
rmse_vals <- rmse_all[-1]                # drop intercept-only
ncomp_seq <- seq_along(rmse_vals)


# Choose optimal number of components by min CV-RMSE
opt_ncomp <- which.min(rmse_vals)
cat("Optimal ncomp by min CV-RMSE:", opt_ncomp, "\n")

```


```
## Optimal ncomp by min CV-RMSE: 6
```

(b) Plotting and choosing numbers of components

```
validationplot(pls_fit, val.type = "RMSEP") # full CV-RMSE curve
```




```
validationplot(pls_fit, val.type = "R2") # CV-R^2 curve
```



```
# Zoom into first 15 components for readability
upper <- min(15, length(rmse_vals))
plot(ncomp_seq[1:upper], rmse_vals[1:upper], type = "b", pch = 19,
      xlab = "Number of components", ylab = "CV RMSE",
      main = "Zoomed CV-RMSE (first components)")
abline(v = opt_ncomp, col = "red", lwd = 2, lty = 2)
points(opt_ncomp, rmse_vals[opt_ncomp], col = "red", pch = 19, cex = 1.4)
text(opt_ncomp, rmse_vals[opt_ncomp], paste0("opt = ", opt_ncomp),
     pos = 4, col = "red")
```

Zoomed CV–RMSE (first components)



Comment

The cross-validated RMSE decreases rapidly as we add the first few components and reaches its minimum at 6 components.


Beyond this point, the error starts to increase again, indicating the beginning of overfitting.

Therefore, the optimal number of PLS components is 6.

(c) Predplot

```
predplot(pls_fit, ncomp = opt_ncomp)
```

y, 6 comps, validation



Comment

The predplot() shows the cross-validated predicted values versus the measured values using 6 components.

Most points lie close to the diagonal, indicating that the model captures the main structure in the data.

There is some scatter, especially at extreme values, but overall the cross-validated fit looks reasonable.

(d) Predicted versus observed values

```
# Prepare test numeric data
test_num <- test[, is_num, drop = FALSE]


# Predict
y_pred <- predict(pls_fit, newdata = test_num, ncomp = opt_ncomp)
y_true <- test_num$y

# RMSE
```

```
rmse <- sqrt(mean((y_true - y_pred)^2))
rmse
```

```
## [1] 0.3638673
```

```
plot(y_true, y_pred,
      xlab = "Observed y",
      ylab = "Predicted y",
      pch = 19,
      main = "PLS: Test Observed vs Predicted")
abline(0, 1, col = "red", lwd = 2)
```



Comment

The PLS model achieves a test RMSE of approximately 0.36, which is higher than the RMSE obtained in the previous exercise (about 0.25). This indicates that, for this dataset, PLS does not outperform the simpler linear model.

A likely reason is that the dataset already has a structured set of lagged predictors and does not suffer heavily from multicollinearity issues, so the advantage of PLS is limited in this case.

The test scatter plot confirms that PLS predictions generally follow the observed values but exhibit more deviation from the diagonal compared to the previous model.

Ex-3


(a) Ridge Regression

```
library(glmnet)

# prepare X and y matrices for glmnet
X_train <- as.matrix(train_num[, setdiff(names(train_num), "y")])
y_train <- train_num$y

# Fit ridge regression (alpha = 0)
ridge_fit <- glmnet(
  X_train, y_train,
  alpha = 0,           # ridge penalty
  standardize = TRUE   # glmnet scales by default
)

# plot the coefficient paths vs lambda
plot(ridge_fit, xvar = "lambda", label = TRUE,
     main = "Ridge Regression Coefficient Paths")
```



Comment

- What does the plot show?

It shows the coefficient paths for ridge regression as the penalty parameter lambda varies.

For large lambda, the coefficients are shrunk strongly towards zero.

As lambda decreases, the coefficients gradually increase in magnitude, approaching the ordinary least squares solution.

- Which default parameters are used for lambda?

glmnet() uses a default logarithmic grid of lambda values (around 100 values), automatically chosen based on data scale.

It starts from a lambda large enough to shrink all coefficients almost to zero and decreases to a small value near the unregularized model.

- What is the meaning of alpha?

alpha = 0 means ridge regression (L2 penalty). Values:

alpha = 0 → Ridge (L2)

alpha = 1 → LASSO (L1)

0 < alpha < 1 → Elastic Net

Ridge shrinks coefficients but does not set them exactly to zero.


(b) Ridge Regression

```
X_train <- as.matrix(train_num[, setdiff(names(train_num), "y")])
y_train <- train_num$y

K <- 10
n_tr <- nrow(X_train)
fold_sizes <- rep(floor(n_tr / K), K)
fold_sizes[1:(n_tr %% K)] <- fold_sizes[1:(n_tr %% K)] + 1
foldid <- inverse.rle(list(lengths = fold_sizes, values = seq_len(K)))

# CV ridge (alpha=0)
set.seed(12321492)
cv_ridge <- cv.glmnet(
  X_train, y_train,
  alpha = 0,                      # ridge
  standardize = TRUE,              # glmnet scales by default
  nfolds = K,                      # will be overridden by foldid, but keep for clarity
  foldid = foldid,                 # consecutive blocks, no time leakage
  type.measure = "mse",            # for Gaussian -> MSE; we'll sqrt later for RMSE
  family = "gaussian"
)
```

```
# Plot CV curve
plot(cv_ridge)
title("Ridge CV - MSE vs log(lambda) (10 consecutive folds)", line = 2.5)
```



```
# Pick lambdas and RMSEs
lam_min <- cv_ridge$lambda.min           # minimizes mean CV error
lam_1se <- cv_ridge$lambda.1se            # 1-SE rule (simpler model)

mse_min <- min(cv_ridge$cvm)             # mean CV MSE at lambda.min
rmse_min <- sqrt(mse_min)

# Also compute RMSE at lambda.1se for reporting
mse_1se <- cv_ridge$cvm[which(cv_ridge$lambda == lam_1se)]
rmse_1se <- sqrt(mse_1se)

cat(sprintf("lambda.min = %.5g; CV RMSE (min) = %.4f\n", lam_min, rmse_min))
```

```
## lambda.min = 0.11958; CV RMSE (min) = 0.2685
```

```
cat(sprintf("lambda.1se = %.5g; CV RMSE (1SE) = %.4f\n", lam_1se, rmse_1se))
```

```
## lambda.1se = 0.33274; CV RMSE (1SE) = 0.2844
```

```
# Coefficients at lambda.min (optional preview)
coef_min <- coef(cv_ridge, s = "lambda.min")
nnz <- sum(coef_min != 0)
cat("Nonzero coefficients at lambda.min:", nnz, "\n")
```

Nonzero coefficients at lambda.min: 108

```
head(as.matrix(coef_min)[,1], 15)
```

	(Intercept)	START.YEAR	START.QUARTER	COMPLETION.YEAR
##	2.308854e+00	6.532715e-03	-1.032039e-02	7.504879e-03
## COMPLETION.QUARTER		PhysFin1	PhysFin2	PhysFin3
##	5.019807e-03	-3.805561e-02	2.198618e-05	2.954967e-05
##	PhysFin4	PhysFin5	PhysFin6	PhysFin7
##	-5.116027e-05	3.325872e-05	1.494950e-04	2.030161e-02
##	PhysFin8	Econ1	Econ2	
##	2.682164e-04	8.919271e-06	8.630951e-05	

Comment

We fitted ridge regression with cv.glmnet using 10 consecutive folds to respect time order.

The CV curve decreases smoothly as lambda decreases and reaches its minimum at lambda_min 0.12, with CV RMSE 0.27.

As a more conservative choice, the 1-SE rule selects lambda_1se 0.33 with a slightly higher CV RMSE (0.28) and stronger shrinkage.

The corresponding ridge coefficients are obtained via coef(cv_ridge, s = "lambda.min") (or "lambda.1se"). Ridge shrinks magnitudes but does not set coefficients exactly to zero.

(c) Ridge — test predictions, plots, RMSE

```
stopifnot(exists("cv_ridge"))

# Ensure test has the same numeric columns as training (after our NZV drop)
test_num <- test[, colnames(train_num), drop = FALSE]

# Matrices for glmnet
X_test <- as.matrix(test_num[, setdiff(colnames(test_num), "y")])
y_true <- test_num$y

# Predict with lambda.min and lambda.1se
y_pred_min <- drop(predict(cv_ridge, newx = X_test, s = "lambda.min"))
y_pred_1se <- drop(predict(cv_ridge, newx = X_test, s = "lambda.1se"))
```

```

# RMSE helper
rmse <- function(y, yp) sqrt(mean((y - yp)^2))

rmse_min <- rmse(y_true, y_pred_min)
rmse_1se <- rmse(y_true, y_pred_1se)

cat(sprintf("Test RMSE (ridge, lambda.min): %.4f\n", rmse_min))

```

Test RMSE (ridge, lambda.min): 0.3251

```
cat(sprintf("Test RMSE (ridge, lambda.1se): %.4f\n", rmse_1se))
```


Test RMSE (ridge, lambda.1se): 0.3404

```


# Plot: observed vs predicted (lambda.min)
plot(y_true, y_pred_min,
      xlab = "Observed y (test)",
      ylab = "Predicted y (ridge, lambda.min)",
      pch = 19, main = "Ridge - Predicted vs Observed (test)")
abline(0, 1, col = "red", lwd = 2)

```

Ridge — Predicted vs Observed (test)



```
# Second plot for 1SE
plot(y_true, y_pred_1se,
      xlab = "Observed y (test)",
      ylab = "Predicted y (ridge, lambda.1se)",
      pch = 19, main = "Ridge (1SE) - Predicted vs Observed (test)")
abline(0, 1, col = "red", lwd = 2)
```



Comment

Using the optimal ridge model, we predicted the response on the test set and compared the results against the observed values.

The scatter plot shows a clear positive linear relationship, with most points lying close to the 45-degree line, indicating good predictive performance.

The test RMSE for the model with `lambda_min` is approximately 0.33, and about 0.34 under the 1-SE rule.

Compared to the previous models, ridge performs better than PLS (0.36) but does not outperform the standard linear regression model from Exercise 3 (0.25).

This is consistent with the idea that ridge improves stability through coefficient shrinkage, but may not produce the lowest error when multicollinearity is moderate and the baseline model already fits well.

(d)

Comment

Even though ridge penalization shrinks the overall L2-norm of the coefficient vector as lambda increases, individual coefficients do not have to move strictly in one direction.

With correlated predictors, the coefficients influence each other, so some paths can bend slightly rather than decrease smoothly.

In addition, glmnet computes the solution path numerically, which can introduce small non-monotonic wiggles.

In other words, the total shrinkage is monotonic, but individual coefficient curves may show small reversals, and this is expected behavior rather than a problem.